

# Portable Spec Kit — Spec-Persistent Development

## Spec-Persistent Development for AI-Assisted Engineering

A single portable file that defines your complete development methodology — personalized from the first session. Drop it into any project and the AI agent detects your GitHub profile, adapts to your expertise and working style, follows your engineering standards, manages tasks, writes tests, and maintains context across sessions.



**Author:** Dr. Aqib Mumtaz, Ph.D.

**Specialization:** Ph.D. Computer Science — Artificial Intelligence


**Research:** Multimodal AI, Healthcare AI, Autonomous Surveillance

**Version:** v0.6.98

**Date:** June 2026

# The Framework

**Spec-Persistent Development** — specs always exist and stay current, but never block. The framework doesn't enforce a methodology — it supports waterfall, agile, or mixed. The AI agent maintains living specifications throughout development, adapting to how you work.



## 11 Core Strengths

| Strength               | What It Means                                                                                                         |
|------------------------|-----------------------------------------------------------------------------------------------------------------------|
| <b>Lightweight</b>     | Single markdown file. Zero dependencies. Zero install.                                                                |
| <b>Portable</b>        | One file → any repo. Works instantly. Symlinks for all agents.                                                        |
| <b>Personalized</b>    | GitHub profile auto-detect. Adapts to your expertise, communication style, and working pattern. Tailored AI behavior. |
| <b>Spec-Persistent</b> | SPECS → PLANS → TASKS → RELEASES. Specs persist alongside your code. Your workflow, your choice.                      |

|                                                                                                             |                                                                                                                                                                                                                                  |
|-------------------------------------------------------------------------------------------------------------|----------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
|  <b>Agent-Agnostic</b>     | Claude, Copilot, Cursor, Windsurf, Cline — one source, all sync.                                                                                                                                                                 |
|  <b>Context-Persistent</b> | AGENT_CONTEXT.md carries state across sessions. Pick up after weeks. Never lose context.                                                                                                                                         |
|  <b>Secure</b>             | Never expose API keys — even if asked. Strict .env rules. Code security enforced.                                                                                                                                                |
|  <b>Non-Blocking</b>       | Code first, specs later. Agent tracks silently. Never blocks the developer.                                                                                                                                                      |
|  <b>Self-Scaffolding</b>   | Auto-creates 9 agent files, README, workspace context. 8 stack templates built-in. Ready to code in seconds.                                                                                                                     |
|  <b>Self-Validating</b>    | Agent tests everything. Fixes before presenting. 98%+ coverage target.                                                                                                                                                           |
|  <b>Reliably Enforced</b>  | 3-layer enforcement: bash sync-check (18 deterministic checks, PSK001–PSK017), sub-agent critic (verbatim-quote gate), git + Claude Code hooks (block bad commits). Agent cannot skip quality gates. Works across all AI agents. |

## Reliability Architecture (v0.5.8+)

The kit guarantees consistency through three independent enforcement layers:

- **Bash critic:** `psk-sync-check.sh` runs 18 deterministic checks PSK001–PSK017 (version consistency across 10+ files, test counts, feature counts, SPECS staleness, R→F→T gate, script permissions, required directories, CHANGELOG/RELEASES content, ARD content, AGENT.md stack, README Latest Release + install list + agent table + flow table, flow doc content, critic prompts meta-check, secrets scanning). Mode auto-detection adapts checks to project type (kit/node/python/go/rust/generic). Runs in under 500ms. Silent on clean. Emergency bypass via `PSK_SYNC_CHECK_DISABLED=1` env variable.
- **Sub-agent critic:** `psk-critic-spawn.sh` uses a file-based protocol to spawn fresh-context sub-agents that verify semantic correctness — flow docs describe current state, release notes match actual changes. Five-iteration cap with human escalation.
- **Hooks:** `.claude/settings.json` PostToolUse warns on every edit if drift detected. `.git/hooks/pre-commit` blocks commits that fail sync-check. These run outside agent control — the agent physically cannot bypass them during commits.  
Install: `curl -fsSL https://raw.githubusercontent.com/aqibmumtaz/portable-spec-kit/main/install.sh | bash` — delivers the complete reliability infrastructure to your

project. Uninstall: `bash agent/scripts/psk-uninstall.sh` (preserves your pipeline files).

## v0.5.23 — Full-Surface Documentation Coverage Analyzer

`agent/scripts/psk-doc-sync.sh` is the release-time tool that catches documentation drift *structurally* rather than by agent inspection. It extracts every feature from the current-minor CHANGELOG entry and checks its presence across five documentation surfaces: `agent/*.md` (architecture), `docs/work-flows/*.md` (user-facing workflows), `docs/research/*.md` (methodology paper), `ard/*.html` (architecture reference), and `README.md` (public overview). Each feature is classified COVERED ( $\geq 2$  surfaces), PARTIAL (1 surface), or MISSING (0 surfaces) with a surface-legend code `[AFPDR]` and a heuristic-suggested target doc per gap.

The analyzer is mandatory at release Step 4 (Update Docs). The `STEP_4_FLOW_DOCS` and `STEP_9_VALIDATION` critic prompts now include an "AUTOMATED COVERAGE TOOL" clause that makes the sub-agent critic run the analyzer and treat every MISSING entry as STALE. Run with `--strict` for CI-style behavior (exit 1 on any gap). This addresses the structural reason v0.5.13–v0.5.18 shipped with flow-doc drift: critic prompts were only catching contradictions, never omissions.

## v0.6.0 — AVACR: Reflex Evolves to Adversarial Goals + Sandbox Isolation + Peer-Exchange

F70 Reflex evolves from cooperative VACR (Verbal Actor-Critic Refinement) to **AVACR — Adversarial Verbal Actor-Critic Refinement Loop**. The loop gains explicit asymmetric goals, physical sandbox isolation, and machine-parseable peer-to-peer YAML exchange so humans read only the final signoff, never agent working output.

- **Adversarial goals with collaborative convergence:** QA-Agent's explicit goal is to FAIL the release (hunt exhaustively across 14 dimensions; default DENIED). Dev-Agent's goal is to FOOLPROOF the release against QA's hunt (fix findings + audit + harden sibling-class weaknesses). Convergence is QA hunting hard and finding zero blockers — not "Dev says it's done."
- **Physical sandbox worktree isolation:** QA operates in a git worktree at `reflex/sandbox/qa-pass-NNN/` with `agent/AGENT.md` and `agent/AGENT_CONTEXT.md` physically removed. Dev's project rules + living session narrative cannot bias QA's fresh-context investigation. Dev retains full-repo access. Dual-rooted access: `REFLEX_QA_WORKTREE` for reads, `REFLEX_QA_PROJECT` for invocations.
- **Peer-to-peer YAML exchange:** QA writes `findings.yaml` (priority + dimension + blast\_radius + confidence + reproducibility + escalation + citable\_quote + regression\_vector + standard + recommendation per finding). Dev returns `dev-`

`result.yaml` (fixed / rebuttals / routed\_to\_human / escalated / deferred). Human reads only `signoff.md` (GRANTED or DENIED with exit criteria).

- **14-dimension adversarial hunt:** functional / edge-case / cross-pipeline / workflow / research-backed-quality / superficiality / production-readiness / security (OWASP) / performance / architectural-compliance / documentation / test-quality / tech-debt / readiness-affirmation.
- **4-persona testing:** every feature exercised as new-user / power-user / malicious-user / accessibility-user — different personas surface different failure classes.
- **Research capability + citable-quote honesty gate:** QA may invoke WebSearch/ WebFetch for RFCs, OWASP Top 10, CVEs, benchmarks, academic papers. Every finding requires a citable quote with file:line reference + verbatim text that bash grep can verify. No hallucinated standards.
- **Dev peer-fixer protocol:** 4-bucket diagnosis (A/B/C/D with Bucket-D requiring counter-citation rebuttal); human-arbitration routing for cross-cutting / ADL-amendment findings (Dev does not auto-fix those); flaky-vs-deterministic distinction; sibling-class hardening mandated (fix the class, not just the instance).
- **Coverage declaration + regression-vector re-execution:** QA writes `coverage.md` listing tested × dimensions AND untested surface with reasons. Regression vectors preserved per-finding and re-executed verbatim each pass; failure upgrades priority one level (catches surface-patch vs root-cause fixes).
- **CHANGELOG truth-check mandate:** every bullet in current version's CHANGELOG must have a test case and observable evidence; unverifiable claims block signoff.
- **Consistency sweep:** `cycle` → `pass` (~30 refs), `Refine` → `Reflex` (3 stragglers with real install/update bugs surfaced and fixed).
- **Paper target:** *"Adversarial Verbal Actor-Critic Refinement (AVACR): Multi-Agent Verbal Reinforcement Learning with Asymmetric Goals for Spec-Grounded Post-Release Program Repair"* (ICSE / FSE / ASE). Note: "Adversarial" refers to inference-time asymmetric goals / red-teaming, not adversarial training (no gradient updates).

### v0.6.1 — Kit self-evolution loop + MINIMAL template tier + 14 framework-gap closures

Dual-track release closing the empirical-kit-evolution loop end-to-end and adding a third response-template tier.

- **N33 adopter-side auto-submit:** consent-gated dispatcher opens anonymized kit issues at pass-end. Three modes (manual / prompt / auto) with four guards (consent, pass filter, 24h rate limit, anonymizer).
- **N38 maintainer-side intake automation:** weekly GitHub Action parses `avacr-eval` issues, drafts `agent/tasks/Gxx-*.md` kit tasks.
- **N44/N45 general-purpose runner:** mode-less `reflex/run.sh [--target <path>] ; --kit-loop` chains prep-release + self-test with max 3 iterations.

- **Per-pass canonical trace (Q1) + trace continuity:** structured `findings.yaml` + verdict `signoff.md` per pass; next QA reads prior `findings.yaml` + `signoff.md` + cross-pass register at Phase 1.
- **MINIMAL response-format tier + Writing Style sub-section + breadcrumb rules 6a/6c/6d + pre-send self-check.**
- **Session-stack 3-tier scalability + promotion rule.**
- **14 framework-gap closures (G1-G15 less G11).**

## **v0.6.2 — Dev-branch isolation + protected-files write-ban + autoloop + history retention**

Closes the reflex execution-model loop. QA was sandboxed; Dev now runs on an isolated branch with a 3-layer write-ban on the two spec-persistent-owned files. Autoloop unifies the iterate-until-GRANTED command for kit + user projects.

- **Dev-branch isolation:** every reflex pass runs Dev on `reflex/dev-cycle-NN-pass-NNN`. Fast-forward merge on GRANTED + auto-delete; unhappy branches retained (last 3) for review.
- **Protected-files write-ban (3-layer):** `agent/AGENT.md` and `agent/AGENT_CONTEXT.md` unmodifiable by reflex — prompt constraint + `gates.sh` diff check + `psk-sync-check.sh` pre-commit hook on reflex dev branches.
- **Sandbox purge after QA:** `file-bugs.sh` removes the current pass's sandbox immediately after findings extraction — Dev physically cannot read QA's workspace.
- **Autoloop generalization:** `bash reflex/run.sh` (default = autoloop) chains prep-release + pass + iterate until GRANTED. Single CLI surface; wrapper scripts retired.
- **Unified cycle-NN-pass-NNN naming:** every pass gets a cycle id (ad-hoc runs auto-assign one); single directory scheme.
- **History retention + reset:** `history_retention` config bounds disk (last 10 pass dirs + 3 dev branches + 3 sandboxes; register + summary.csv forever). `bash reflex/run.sh --reset [--confirm]` for full wipe.
- **Concurrency + empty-pass + intake parser update.**

## **v0.6.3 — Autonomous reflex run validation**

Mechanical patch minted from an autonomous end-to-end reflex run on the v0.6.2 consolidation. Prep-release validated the full reflex architecture through the kit's own prepare-release pipeline — Dev-branch isolation, protected-files write-ban, autoloop, history retention, reset command, unified cycle naming all verified green. No feature changes.

## v0.6.4 — Reflex convergence stopping + loop-resume fix + gates cache fix + token tracking

- **Convergence-based stopping (N56)** replaces "max 3 iterations" with real signals: GRANTED, REGRESSION, findings\_floor, plateau ( `patience_passes` ), fix-rate drop ( `min_fix_rate` ), safety cap.
- `--loop --resume` **collision fix**: deferred flag resolution in `run.sh` ; `--loop --resume` now routes correctly to `loop-resume` mode instead of hijacking to `resume-qa` .
- `gates.sh` **cache-order bug**: rft-cache cleared at gate-runner entry so sequential gates don't see stale data.
- **Token tracking ( `reflex/lib/track-tokens.sh` )**: per-pass tokens aggregated into `reflex/history/token-usage.csv` with `tokens_per_finding` metric + per-pass / per-cycle budget warnings.
- **Iteration auto-chain on gate-fail**: spurious gate failures no longer abort the autoloop — only REGRESSION + verdict signals halt.

## v0.6.5 — Kit-bootstrap integrity gate (4-layer defense against un-installed-kit projects)

Closes the failure mode where a non-kit-aware agent (Copilot, plain-Claude) scaffolds a project manually, commits a `v0.N.N:` subject that satisfies reflex Gate 2, and proceeds end-to-end through prep-release + reflex — all while the project has zero kit scripts, no hooks, no skills, no CLAUDE.md, no `.portable-spec-kit/config.md` . Surfaced by an audit of the `psk-playground/urdu-stt-distillation` project which passed 3 reflex GRANTED verdicts despite being infrastructurally empty.

- **Layer 1** — `install.sh` (unchanged, canonical installer).
- **Layer 2** — `psk-release.sh` **Step 0**: `run_bootstrap_gate()` fires before `init_state` on `prepare` / `refresh` . FAIL exits 1 with full bootstrap report + remediation command — no release state written.
- **Layer 3** — `reflex/lib/preconditions.sh` **Gate 0a**: runs after Gate 0 (self-test identity) and before Gate 1 (clean tree). Script resolution prefers target's own copy, falls back to kit's copy — catches the exact case where target has zero scripts.
- **Layer 4** — **QA-Agent Dimension 16 "Kit-bootstrap integrity"**: new mandatory dimension in `reflex/prompts/qa-agent.md` . Runs FIRST (before Phase 1 Plan). Halts pass with `QA-BOOTSTRAP-NN CRITICAL` finding on any C1-C7 fail.
- **New** `agent/scripts/psk-bootstrap-check.sh` (298 lines, shared helper). 7 checks (C1-C7): framework file + entry-point symlink, `.portable-spec-kit/config.md` , core 6 kit scripts executable,  $\geq 10$  cached skill files, pre-commit hook wired to `psk-sync-check`, all 9 `agent/*.md` pipeline files, test harness. Modes: default report, `--quiet` , `--json` , `--remediate` (auto-invokes parent `install.sh`). Kit-self auto-detected — skips symlink/config checks for the kit repo.



- **Register format improvements** in `reflex/lib/update-eval-trace.sh` : parses cycle from pass-name when `.cycle-meta` has stale `cycle=0` , adds `coverage.md` drill-down, extracts "Tests exercised:" summary line per pass (regression vectors + dimensions).  
Bypass for genuine emergencies: `PSK_BOOTSTRAP_CHECK_DISABLED=1` . Impact: structurally impossible now for an un-bootstrapped project to run prep-release or reflex.

## **v0.6.6 — Reflex Phase 0 pre-compute (claim-driven probing + reference-state diff + assumption surfacing)**

Closes the meta-gap surfaced in v0.6.5 review: reflex's QA-Agent ran 3 GRANTED passes on an infrastructurally-empty project yet missed every kit-infrastructure gap because its probe set was a closed 16-dimension checklist from the AVACR paper. Adding dimensions reactively (Dim-15, Dim-16) fixes one case each but does not prevent the NEXT unknown class — QA always misses the bug outside its list. v0.6.6 moves reflex from "fixed checklist" to "probe-set derived from the project itself," **while keeping the 16-dim checklist as a safety net**. Three modes compose (each deterministic bash pre-compute — <1s, <100 tokens):

- **Mode A — Claim-driven probing ( `reflex/lib/extract-claims.sh` )**: walks every public doc (README, portable-spec-kit.md, SPECS, CHANGELOG, RELEASES, AGENT\_CONTEXT, qa-agent.md) and extracts every claim with `probe_type` + `probe_target` . QA verifies each or files a finding. Unverified or vague claim = finding. Probe types: version-match, test-count, feature-count, capability-exists, file-count, install-works, dimension-count.
- **Mode B — Reference-state comparison ( `reflex/lib/state-diff.sh` + `reflex/reference-state/speckit-project.yaml` )**: static definition of what a complete speckit project must contain (required files + pipeline files + dirs + git hooks + entry-point symlinks + retired-file exclusions). State-diff compares actual vs reference; deltas pre-classified by severity; CRITICAL deltas auto-promote to findings. Kit-self auto-detected — skips user-project-only checks.
- **Mode C — Assumption surfacing (in `reflex/prompts/qa-agent.md` )**: every pass writes `assumptions.md` listing implicit assumptions ("kit is installed," "test harness exits 0 on green," "every [x] feature actually implemented"). Each assumption gets a probe. Unverifiable assumption = MAJOR `QA-ASSUMPTION-NN` finding.
- **QA Phase 0 wiring**: both helpers run in `reflex/lib/spawn-qa.sh` BEFORE sandbox creation so `claims.yaml` + `state-diff.yaml` land in the pass directory pre-populated. QA reads these in Phase 0 — no LLM cost to re-derive. The 16-dim checklist runs as a safety net. CRITICAL state-diff short-circuits the 16-dim sweep of a structurally-broken project.



- **Dev-Agent expanded scope (in `reflex/prompts/dev-agent.md`):** Dev closes three gap classes per pass — (a) findings from `findings.yaml`, (b) unverified claims from `claims.yaml`, (c) state-diff deltas from `state-diff.yaml`. "Build, don't just patch" — restore missing infrastructure, build unfulfilled capabilities, add structure that makes assumptions self-verifying. Scope guardrails preserved (no out-of-SPECS features, protected files banned).

**Design discipline:** extension within AVACR/Reflex, NOT a new framework. QA is still Critic, Dev is still Actor, adversarial-with-collaborative-outcome loop unchanged. What changed: QA's probe-set derivation (paper-defined → project-derived) and Dev's scope (patch-only → patch + build + remediate).

## v0.6.7 — Reflex 7-layer Senior/Principal-level QA system

The user must NOT be QA of their own project. Reflex now operates with a 7-layer Senior/Principal-level QA system that derives all probes from the project's spec pipeline + kit toolkit + running system — zero human-curated lists. Builds on v0.6.6 Phase 0 pre-compute by adding four new verification layers and codifying Senior/Principal philosophies in the prompts.

- **Layer 1 — R→F→T pipeline integrity** ( `reflex/lib/check-rft-integrity.sh`, deterministic bash, Phase 0). For every `[x]` feature in SPECS.md: maps to Rn in REQS.md (R→F traceability), has acceptance criteria block, has `agent/design/f{N}-*.md` design plan, Tests column references real test file, test contains non-trivial assertions (no `toBeDefined` trivia / `test.skip` / TODO markers), test exits 0, has `[x]` mark in TASKS.md. Each break = CRITICAL finding. Output: `rft-integrity.yaml`.
- **Layer 2 — Bidirectional doc ↔ code consistency** ( `reflex/lib/doc-code-diff.sh`, deterministic bash, Phase 0). Doc-to-code: every named file/script/version/count/capability claim → grep code for backing. Code-to-doc: every kit script + every `[x]` feature → at least one doc surface mentions it. Numerical-claim cross-check (test count consistency). Drift = finding. Output: `doc-code-diff.yaml`.
- **Layer 3 — Behavioral per-feature verification** (qa-agent.md Phase 2 mandate, LLM). Per `[x]` feature: 1 happy + 5 edge cases + 3 adversarial inputs + 1 integration scenario. **Independence rule:** design own test BEFORE reading Dev's. Output: `behavioral-tests/F{N}-*`.
- **Layer 4 — Test-quality audit** (qa-agent.md Phase 3 mandate, LLM). After Layer 3, audit Dev's test: would it catch QA's adversarial inputs? Are mocks honest? Coverage adequate? Inadequate Dev test = MAJOR `QA-TEST-INADEQUATE-F{N}` even if currently passing. Output: `test-quality-audit.yaml`.
- **Layer 5 — Cross-feature integration probes** (qa-agent.md Phase 2 mandate). Identifies pairs of `[x]` features that interact (auth × data, search × permissions).

Surfaces compose-time bugs single-feature tests miss. Output: `integration-probes.md` .

- **Layer 6 — Production readiness via kit tools** ( `reflex/lib/spawn-qa.sh` Phase 0 capture). Invokes `psk-sync-check.sh --full` , `psk-doc-sync.sh` , `tests/test-release-check.sh agent/SPECS.md` , `psk-code-review.sh` . Every PSK error code, MISSING in doc-sync, R→F→T break in release-check becomes evidence. Output: `kit-tools-output/{sync-check,doc-sync,release-check,code-review}.txt` .
- **Layer 7 — External reality check** (`qa-agent.md` Phase 2 mandate, LLM + WebSearch/WebFetch). For code touching network/dependencies/standards: cross-check OWASP Top 10, CVE feeds, framework deprecations (e.g., transformers `as_target_processor()` removed in v4.44), best-practice patterns. Output: `external-research.md` .
- **Senior/Principal-level philosophies** codified in `qa-agent.md` ("the user MUST NOT be QA of their own project") and `dev-agent.md` ("Senior/Principal engineer who closes gaps QA surfaces and hardens implementation against future passes"). 7 operating principles each, 4 forbidden moves, three gap classes per pass for Dev (fix findings + build unfulfilled claims + remediate state-diff deltas).
- **Architecture plan committed:** `agent/design/f70-reflex-senior-engineer-qa.md` documents context, philosophies, all 7 layers, what stays from v0.6.5/v0.6.6, what was discarded, v0.6.7-v0.6.13 sequencing. Companion Dev plan at `agent/design/f70-reflex-senior-engineer-dev.md` . Layer 4 hardening log shipped ( `reflex/history/hardening-log.md` ).

**Short-circuit rule:** if Layer 1 (rft-integrity) or Mode B (state-diff) shows CRITICAL deltas indicating structural break, halt the pass with those findings and skip the 16-dim sweep — fix structure first, behaviors second.

## v0.6.8 — Reflex 7-layer full implementation: Layer 3/5/7 QA seed-helpers + Layer 7 Dev ADL extraction

Completes v0.6.7's architecture by shipping deterministic helpers for Layers 3, 5, 7 (QA) and Layer 7 (Dev). All 7 layers now have deterministic-bash seed-helpers wherever possible — LLM creative work goes to genuine reasoning, not derivation.

- `reflex/lib/scaffold-behavioral-tests.sh` (**QA Layer 3**): for every `[x]` feature in `agent/SPECS.md` , generates `behavioral-tests/F{N}-test-plan.md` skeleton with the 1+5+3+1 structure (1 happy + 5 edge + 3 adversarial + 1 integration). Pulls acceptance criteria from `### F{N}` blocks; flags missing criteria as Layer 1 finding. QA fills TODOs with concrete probes — independence rule: design BEFORE reading Dev's test. Skeleton saves ~5K tokens/pass.
- `reflex/lib/identify-integration-probes.sh` (**QA Layer 5**): scans SPECS for feature interactions via design-plan + criteria-block cross-references. Emits `integration-`

`probes.md` with candidate pairs. QA validates pairs, drops spurious, designs scenarios. Includes "common integration patterns" reference.

- `reflex/lib/external-research.sh` (**QA Layer 7**): scans project manifests (`package.json`, `requirements.txt`, `pyproject.toml`, `go.mod`, `Cargo.toml`, `Dockerfile`) and emits `external-research.md` with seeded WebSearch/WebFetch queries: per-dependency CVE searches, framework-changelog URLs, OWASP Top 10 mapping for project class. Stack-specific recent-vulnerability-class hints (e.g., prompt injection per OWASP LLM Top 10).
- `reflex/lib/auto-extract-adl.sh` (**Dev Layer 7**): post-pass, scans Dev's commits for substantive bodies ( $\geq 200$  chars) containing rationale keywords. Drafts ADL entries to `agent/.release-state/adl-drafts.md` for kit-maintainer review.
- **Phase 0 produces 8 artifacts**: `spawn-qa.sh` invokes all helpers BEFORE sandbox creation. Pass dir contains `claims.yaml`, `state-diff.yaml`, `rft-integrity.yaml`, `doc-code-diff.yaml`, `behavioral-tests/`, `integration-probes.md`, `external-research.md`, `kit-tools-output/` when QA spawns.
- **Field-tested on kit-self**: Layer 3 produces 70 test-plan skeletons in  $< 1s$  · Layer 5 finds 77 candidate pairs · Layer 7 QA gracefully detects stack signals · Layer 7 Dev drafts ADL entries when commits have rationale.

### v0.6.9 — The kit is audited and improved by the reflex it built

The architecture built across v0.6.5-v0.6.8 — Phase 0 deterministic helpers + Senior/Principal-level QA philosophy + paired Dev operating system — was designed to find drift in user projects. v0.6.9 turned that exact machinery on the kit itself. The same QA-Agent that catches bugs in your code caught real bugs in the code that builds the QA-Agent.

- **Phase 0 helpers ran on kit-self** and exposed **180 mechanical drift items** the kit's own existing mechanical gates had passed clean for months: 170 R→F→T pipeline breaks (55 features missing `### F{N}` blocks, 45 missing design plans) + 10 doc↔code drifts including 1 LIVE rename incompleteness from v0.6.2 entry consolidation.
- **QA-Agent was spawned** against the kit and filed **5 findings** (3 MAJOR + 2 MINOR), prioritized for the field-test scope.
- **Dev-Agent closed 2 inline**: `QA-DOC-FT-02` live `kit-loop-state.yml` → `loop-state.yml` rename (4 files); `QA-TEST-INADEQUATE-F1` Layer-4 hardening shipping new `PSK018` per-feature pairwise check (advisory by default; `PSK_FEATURE_CRITERIA_STRICT=1` promotes to hard fail). The PSK018 hardening would have caught a "69 of 70 acceptance-criteria blocks deleted" silent ship under the existing shallow gate.
- **Dev deferred 3 to v0.6.13 ARB** in `agent/TASKS.md`: `QA-PIPE-FT-01-ARB` (backfill 55 missing criteria blocks — kit's own RFT debt) + `QA-INT-F70-F52-ARB` + `QA-INT-`

F70-F1-ARB (behavioral integration probes need full QA+Dev pass, not field-test scope).

- **Layer 5 hub-suppression added** in `reflex/lib/identify-integration-probes.sh` : features appearing in  $\geq 10$  candidate pairs are hubs (e.g. F64 design pipeline cross-refs everything; F70 reflex links many) — pairs involving hubs are mostly noise, suppressed into a separate callout. Result on kit-self: 77  $\rightarrow$  13 pairs.
- **Field-test artifacts committed** in `reflex/history/field-test-v0.6.8/` as permanent evidence: claims, state-diff, rft-integrity, doc-code-diff, behavioral-tests/, integration-probes, external-research, kit-tools-output, findings, signoff, dev-trace.
- **Hardening event H-001 logged** in `reflex/history/hardening-log.md` .

*A tool sufficient to audit user projects must be sufficient to audit the project that built it — v0.6.9 proved that proposition. v0.6.13 then closed the kit's own RFT debt + flipped PSK018 strict, locking in the standard the architecture proved.*

## **v0.6.13 — Nested per-cycle history layout + Dim 23 + cycle-summary aggregator + autoloop convergence**

Closes the meta-gap class user surfaced in cycle-01 trace audit (5 issues + 4 sub-gaps), formalizes Dim 23 so QA self-evolves to catch the class going forward, migrates reflex history to a nested per-cycle directory layout, and drops the hard-3-iter autoloop cap in favor of convergence-based stopping.

- **Dim 23 — Auditor-output hygiene (MANDATORY):** five new probes (register hygiene · cost data persistence · parallel-directory disambiguation · cross-surface terminology consistency · self-output validation loop) catch the meta-gap class previous 22 dimensions didn't probe. `reflex/prompts/qa-agent.md` v0.6.13 grows from 22  $\rightarrow$  23 dimensions; QA's own self-evolution loop (Dim 18) is now the mechanism by which new dimensions are added.
- **Nested per-cycle history layout:** `reflex/history/cycle-NN-pass-NNN/`  $\rightarrow$  `cycle-NN/pass-NNN/` (autoloop) and `standalone/pass-NNN/` (single-pass). Sandbox mirrors history layout. Per-cycle parent dir hosts a `summary.md` aggregator ( `reflex/lib/cycle-summary.sh` ). Flat identity ( `cycle-NN-pass-NNN` ) reserved for git branches, CSV pass id, and `[source: ...]` commit trailers — decouples branch/identity from on-disk layout.
- **Autoloop runs until convergence:** hard 3-iter cap dropped. Primary stops are convergence-based (GRANTED · REGRESSION · findings\_floor · plateau · fix-rate drop). `convergence.max_iterations_safety` default 20 is escape hatch only.
- **Cycle-01 trace audit fixes (9/9):** pass-002 visible in register as INCOMPLETE (was silently dropped); "Autoloop cycle"  $\rightarrow$  "Reflex cycle" terminology unified; sandbox-vs-history disambiguation in `17-reflex.md` ; `score.sh` `grep -c || echo 0` multi-line bug fixed ( `head -1` ); `doc-code-diff.sh` path-with-spaces fix (10 $\times$  scan coverage 12  $\rightarrow$  120 docs); `psk-doc-sync.sh --strict` tightened (Missing>0 only).

- **First successful Dev-Agent run end-to-end:** `cycle-01-pass-004` closed `QA-NOISE-RE-01` bucket A (NOISE\_RE backtick-prefix SHA gap), 1 commit, all 6 mechanical gates green, fast-forward auto-merge to main — first time QA→Dev→merge loop closed without manual intervention.
- **3-layer QA→Dev contract enforcement:** `resume-qa` requires `findings.yaml` + `signoff.md` ; `resume-dev` requires `dev-trace.md` ; `loop.sh` Phase 4 INCOMPLETE detection. Closes "Dev never spawned on partial QA timeout" silent-skip pattern observed in `cycle-01-pass-002` .
- **Tests:** 1633 passing (1488 framework + 145 benchmarking). Sync-check 19/19. Release-check 70/70 features pass with R→F→T coverage.
- **Tagline:** *Auditor-output hygiene formalized as Dim 23; reflex history reshaped per-cycle; autoloop runs to convergence.*

## v0.6.12 — Reflex iter-1 convergence anchor + close all 15 v0.6.11 audit findings

Anchor version for the reflex iter-1 convergence cycle. Consolidates 8 refresh-release commits from v0.6.11 baseline. All 15 audit findings (5 pre-existing ARBs + 10 iter-1 QA findings) resolved or properly deferred; kit's own QA infrastructure self-validated end-to-end via reflex iter-1 self-test on kit-self.

- **All 10 iter-1 QA findings closed:** QA-REL-NONDETERM-01 (CRITICAL — self-cleaning rft-cache), QA-KIT-RFT-01 (anchored regex, 70 false-positives → 0), QA-KIT-DCD-02 (DOC\_SURFACES expanded 6→11), QA-PERF-PHASE0-01 (persistent test-ref cache, 10.9× speedup), plus 5 MINOR + Option C thematic test-suite split (closes QA-TEST-COUPPLING-01).
- **5 pre-existing ARBs closed:** QA-ARCH-01 (Chrome → WeasyPrint per ADR-006), G2/G5 (verified already-shipped), QA-INT-F70-F1-ARB (C1.5 symlink target validation), QA-INT-F70-F52-ARB (layered R→F→T documented as intentional).
- **Cost/perf wins:** persistent test-ref cache cold 54s → warm 5s = 10.9× speedup. Phase 0 wall-clock 74s → ~25s on subsequent invocations. Dim 17 cost/perf audit + Dim 18 philosophy self-audit added to QA prompt + Dev mirror principles 8/9. `pass_score` (Dim 17 ranking) v3 summary.csv schema.
- **Tests:** 1621 passing.
- **Tagline:** *Iter-1 self-audit complete — kit closed every finding it surfaced about itself.*

## v0.6.11 — Reset allowlist + iter-aware refresh-release optimization

Two kit-infrastructure improvements driven by usage friction noticed during v0.6.10 wrap-up. Both reduce maintenance burden + cost without changing reflex semantics.

- **Reset reflex — allowlist-based collection:** `reflex/lib/reset.sh` switched from hardcoded glob lists to allowlist-based `find -mindepth 1 -maxdepth 1` traversal. New

`HISTORY_KEEP=("hardening-log.md")` allowlist preserves kit's structural-defense audit memory across resets. Symmetric `--reset-hardening` flag for true clean-slate.

- **Iter-aware release-prep:** `reflex/lib/loop.sh` Phase 1 branches on `ITER`: iter 1 = `prepare` (version bump), iter 2+ = `refresh` (no bump). Single converged reflex cycle no longer inflates versions.
- **Doc-coverage backfill:** new "v0.6 Capability Reference" section in `docs/work-flows/17-reflex.md`. Closes 24 cumulative `psk-doc-sync.sh` coverage gaps from v0.6.5-v0.6.10.

## v0.6.10 — Closing v0.6.9 ARBs: kit's own RFT debt backfilled + PSK018 strict by default

v0.6.9's field-test surfaced that the kit itself had accumulated RFT (Requirements → Features → Tests) debt: 55 of 70 features lacked their per-feature `### F{N}` acceptance criteria block in `agent/SPECS.md`. The new Layer-4 hardening shipped in v0.6.9 ( `PSK018` ) caught it as advisory; v0.6.10 acted on it.

- **Backfilled all 55 missing `### F{N}` acceptance criteria blocks** in `agent/SPECS.md` (F1–F55, 3-5 testable bullets each, ~200 lines), closing `QA-PIPE-FT-01-ARB`. F56–F70 blocks already existed and were preserved unchanged.
- **Flipped `PSK018` advisory → strict-by-default** in `agent/scripts/psk-sync-check.sh`. Default: `local strict="$ {PSK_FEATURE_CRITERIA_STRICT:-1}"`. Future commits adding an `[x]` feature without matching `### F{N}` block now fail pre-commit. Opt-out via `PSK_FEATURE_CRITERIA_STRICT=0` for genuine emergencies.
- **Layer 1 R→F→T integrity on kit-self:** 55 C2-criteria breaks → 0. Total drift items 170 → 110 (remaining are C1 R→F + C3 design-plan + minor edges, addressable in future versions).
- **Tests:** 1604 passing (1459 framework + 145 benchmarking). Sync-check 19/19. Release-check 70/70 features pass with R→F→T coverage.
- **Still deferred (v0.6.13+ ARB):** `QA-INT-F70-F52-ARB` (reflex × test-harness behavioral integration probe) + `QA-INT-F70-F1-ARB` (reflex × framework-file behavioral integration probe). Both require full QA + Dev pass scope — not field-test scope. Tracked in `agent/TASKS.md`.

## v0.6.41 — PKFL + convergence discipline + ARD user guide + mandate audit extensions

- **Convergence discipline fix (critical):** Rule 4 in `loop.sh` + `run.sh` — DENIED + 0 unclosed findings now stays in the same cycle (forces QA re-verification); was falsely advancing the cycle counter. `resume-dev` now checks QA signoff before writing GRANTED verdict.



- **PKFL — Project-to-Kit Feedback Loop:** `reflex/lib/kit-evolution.sh` (6-phase orchestrator: discover → G3 no-dead-letter → kit Dev-Agent pass → G5 test suite → G6 smoke → cascade), `check-kit-genericity.sh` (G2 vocabulary scanner: blocks domain nouns / hardcoded routes in kit commits), `smoke-test-examples.sh` (G6 multi-project install smoke test). Gates 10 (G1 file-scope) + 11 (G2 genericity-proof) added to `gates.sh`, fire on `REFLEX_KIT_EVOLUTION=1`. G4 `genericity_proof` field required before routing to kit in `file-bugs.sh`. Flow doc `22-project-kit-feedback-loop.md` created. §PKFL mandatory rule section added to `portable-spec-kit.md`; gates count 9 → 11.
- **ARD user guide + CI generation:** `psk-generate-user-guide.sh` (generates `ard/user-guide.html` from README + AGENT.md), `psk-generate-ci.sh` (auto-creates `.github/workflows/ci.yml` from stack detection). Both wired into `psk-release.sh` Steps 1 + 7 as non-fatal hooks.
- **Dim 25 mandate audit extensions:** Three new checks in `mandate-audit.sh`: ARD user guide absent (ADVISORY), DB migration scaffold absent for Drizzle/Prisma projects (MINOR), shadcn/ui components empty (MINOR). Stack-specific fix patterns (drizzle migrations, shadcn/ui, REQS status sync, UI test quality) added to `dev-agent.md`.
- **Tests:** 1837 framework + 145 benchmarking = 1982. N62/D updated to expect stay-in-cycle behavior per convergence discipline fix.

### v0.6.98 — Reflex cycle-03 convergence: Dev-Agent fixes 10 scope:kit findings + monitor auto-render + redispatch phantom-spawn fix

- The reflex cycle-03 kit-self-test convergence pass: the QA wave surfaced 14 actionable findings (5 MAJOR); the Dev-Agent fixed 10 `scope:kit` findings end-to-end under kit-evolution mode (1 `scope:meta` routed, 3 advisory deferred); all 15 mechanical gates passed green.
- **KIT-GAP-0142** (`reflex/run.sh`): the `--resume-dims` detect loop stalled reporting phantom missing dim-sets — the spawn-parsing awk matched nested `prior_findings` ids as phantom records; fixed by skipping any record lacking both `artifact` and `assigned_dims`.
- **KIT-GAP-0143** (`reflex/lib/progress-surface.sh`): the QA/Dev monitor tables were agent-hand-built (off-by-one risk); now the ONE shared emitter auto-renders the accurate `status --table` from the kit, driving the tracker deterministically and filling each result from the pass artifacts.
- Cycle-03 Dev fixes added a new **PSK050** prose-claim drift gate plus 9 more `scope:kit` fixes; flow-doc currency updates (15 gates, 3020/3165 test counts, PSK001-PSK050 range). 3165 tests passing, 0 failed.



## v0.6.97 — Reflex cycle-03 kit-self-test: generic single-sourced progress-table + structural chunked-drive guard + cycle-summary empty-cycle fix

- **KIT-GAP-0139 — cycle-summary skips bare just-advanced cycle dirs:** `reflex/lib/cycle-summary.sh` `discover_cycles` now yields only cycles with  $\geq 1$  pass dir. A `reflex --next-cycle advance` pins `.active-cycle=N` before cycle-N's first pass exists, leaving a bare `cycle-NN/` dir; `discover_cycles` listed it and `summarize_cycle` returned 1, failing the aggregate run (test 98.10). Now bare just-advanced cycles are skipped (informational); the explicit-`$1` path still returns non-zero for a named empty cycle.
- **KIT-GAP-0140 — psk-tests-gate structural no-inline-fallback guard:** `run` previously ran the full suite inline when no pre-verify marker existed — including in an agent-driven context, where one opaque background block surfaces zero per-section progress in the turn-based chat. Now `run` REFUSES (exit 7 + chunk-drive recipe) in an agent context (no TTY, not CI, no `PSK_TESTS_GATE_ALLOW_INLINE=1`), forcing the chunked path; CI/tty/opt-in still run inline. Same no-inline-fallback pattern as `psk-spawn.sh`. Regression: 109.10 + 109.11.
- **KIT-GAP-0141 — generic single-sourced progress-table:** each chunk carries a static `# stage:<name>` bash-comment label; the ONE central renderer ( `psk-chunked-run.sh` ) counts distinct stages and shows the `Stage` column only when  $\geq 2$  ( `all-tests` → test-spec-kit/benchmarking/release-check), suppressing it for a single-suite drive. `psk-spawn.sh` 's stale hand-written column lists now reference the canonical `status --table`; new sync-check **PSK049** guards single-source; the No-Silent-Wait recipe mandates pasting the table into the reply message body. Regression: 107.18 (+ 107.17 live-stage path preserved).

## v0.6.96 — Kit-evolution Dev-fix path: the reflex Dev-Agent fixes scope:kit findings end-to-end

- **KIT-GAP-0129 — kit-evolution.sh pause-aware:** Phase 3 ran `if ! bash run.sh single` and treated `run.sh` 's exit-2 (AWAITING the multi-author QA dim-agent spawn) as a FAILURE, aborting every kit-evolution pass at preflight-end. Now exit 2/3 is a pause the main agent drives ( `--resume-dims / --resume / --resume-dev` ); the wrapper records pause state, propagates exit 2, and re-entry idempotent-skips a completed pass.
- **KIT-GAP-0130 — Dev-Agent kit-fix authorization:** `REFLEX_KIT_EVOLUTION=1` only activated the genericity gates before; it did not tell the Dev-Agent it could FIX `scope:kit` findings (default routes them to maintainer review). Now `spawn-dev.sh` injects a per-pass authorization gated strictly on the env var; `$Reflex-Finding-Classification` carries the matching exception so the prompt and the kit rule agree. `scope:meta` stays maintainer-routed.

- **KIT-GAP-0131 — G1 file-scope kit-prefix check:** the enumerated-subdir allowlist false-flagged the kit's own files ( `ard/` , `agent/tasks/` , `reflex/history/` ) when the kit is nested in a parent repo, and `--relative` hid sibling bleed. Redesigned to flag only paths OUTSIDE the kit prefix ( `git rev-parse --show-prefix` ) — robust in both topologies, catches real sibling-project bleed.
- **Dev-Agent fixes (12 scope:kit findings):** `jq`→`python3` JSON merge (removes the `jq` dependency), `set -eo pipefail` hardening, 7 previously-untested scripts now covered ( `tests/sections/98-script-coverage.sh` ), prose-constant + longitudinal-drift audit-probe precision, `install.sh` lib coverage, and the `curl-install OWASP-A08` accepted-risk doc. 4 net-new-feature findings routed to non-blocking backlog (G43-G46). 3165 tests passing (3020 framework + 145 benchmarking), 0 failed.

## **v0.6.95 — Fail-closed tests pre-verify gate: in-chat progress for reflex/prepare**

- **KIT-GAP-0123 — fail-closed tests pre-verify gate ( `psk-tests-gate.sh` , new):** the long test suite is decoupled from the release ceremony's opaque `step-1-tests` . The agent drives the suite CHUNKED first (per-section chat progress via `psk-chunked-run.sh` ); each test process writes its own real exit-0 STAMP; on all-green the gate SEALS a fail-closed pre-verify marker keyed to a full tree fingerprint ( `HEAD` + working tree). `step-1-tests` routes through `psk-tests-gate.sh run` , which SKIPS the inline re-run only when the marker proves the current tree — so reflex/prepare surfaces the suite per-section in the chat and step-1 is a no-op (no double-run). Any code/test change, dirty tree, missing/failed unit, or non-git → tests run. Marker + stamps gitignored.
- **KIT-GAP-0123-HARDENED — adversarial-review redesign:** a 4-lens review of the first HEAD-SHA design found 3 CRITICALs (dirty-tree-at-same-HEAD shipping untested edits, the seal grepping agent free-text instead of real test output, zero-results sealing green). Redesigned fail-closed and proof-based: the marker seals only from real exit-0 stamps, every stamp + marker bind the tree fingerprint, non-git refuses to seal AND to skip, and the fail-open mtime window was removed.
- **KIT-GAP-0124 — gate-state isolation:** `PSK_TESTS_GATE_STATE_DIR` points section-109's gate exercise at a throwaway dir so it can't clobber production stamps during a live chunked-drive.
- **PSK048** enforces the chunked-suite pre-verify wiring (gate present + step-1 routed + reflex Phase-1 chunk-driven). KIT-GAP-0116 wires reflex QA/Dev progress through the centralized `psk-chunked-run.sh` . Tests 109.1-109.9; flow doc `34-no-silent-wait.md` updated. 3165 tests passing (3020 framework + 145 benchmarking).

## **v0.6.94 — No-silent-wait centralized progress: chunked-drive + status --table + whole-suite coverage**

- **KIT-GAP-0113 — centralized chunked-drive ( `psk-chunked-run.sh` , new):** chat messages are the only progress surface every client renders (the Claude Code `statusLine` bar is CLI-terminal-only — the VS Code extension / app UI does not render it), so chunked-drive is now the DEFAULT for every long op: it runs the op as a series of short background tasks, and each chunk-completion becomes a chat turn that prints the result + launches the next chunk. Registry via `--suite test-spec-kit / all-tests / prepare-release ; status --table` renders the canonical `| Chunk | Unit | Result |` progress table.
- **KIT-GAP-0114 — whole-suite coverage:** the `--suite test-spec-kit` derivation chunked only sections, silently omitting the 85 feature files (~430 tests); fixed to cover the whole suite.
- **KIT-GAP-0109/0110/0111 — `psk-progress.sh` hardening:** `--statusline` surface, TMPDIR-independent live dir, instant live-file write at op start.
- **PSK047** now also requires the chunk driver present; new flow doc `34-no-silent-wait.md` ; §No-Silent-Wait + §Session Health framing flipped so the breadcrumb (chat) is the guarantee and the `statusLine` is a CLI-only complement. 3165 tests passing (3020 framework + 145 benchmarking).

## **v0.6.93 — Session-health indicator: structural statusLine + UserPromptSubmit loop fix + authoritative ctx limit**

- **KIT-GAP-0097 — `ctx:` monitor surfaces structurally:** migrated the hook event `Stop → UserPromptSubmit` (a `Stop` hook returning `additionalContext` re-invokes the agent every turn, so the badge looped); new `--statusline` mode renders `ctx:` (+ `opt:` ) to the Claude Code status bar every turn independent of the agent. `resolve_limit` derives the context window authoritatively from a `compact_boundary` 's `compactMetadata.preTokens` instead of a 220k guess.

## **v0.6.92 — Multi-author reflex enforcement: structural inline-prevention + detect-and-redispatch**

- **KIT-GAP-0098 — inline-prevention keystone:** `reflex/run.sh` 's orchestrator may write only planning artifacts; `findings.yaml` is synthesized solely from `partial-findings-dims-*.yaml` emitted by real dim-agents, so a multi-author QA pass can no longer silently collapse to inline synthesis. Dim-agent count is now dynamic ( `ceil(dims / max_dims_per_spawn)` ). Detect-and-redispatch resolves qa-prefixed spawn ids via the wave manifest's `artifact:` path.

## v0.6.91 — Cycle-01 clean-restart + kit-fidelity state-machine fixes

- **KIT-GAP-0093 — psk-release.sh staleness guard repointed:** `state_is_stale` / `run_next` / `show_status` were reading the vestigial `agent/.release-state/state` file; repointed to canonical `agent/.workflow-state/release.state`. The legacy file isn't refreshed by `prepare`, so a stale sentinel `RUN_ID` made `psk-release.sh next` falsely refuse with "stale state — version drift risk" on a fresh run.
- **Flow-doc currency:** `docs/work-flows/13-release-workflow.md` Flow-Diagram + Step-9 boxes read "38 checks" vs the table/prose "49"; aligned all five references to 49 (step-4 flow-doc critic).
- **KIT-GAP-0092 logged:** `psk-release.sh` advertises a `reset` verb that `psk-dispatch.sh` rejects ("Unknown verb: reset").
- **KIT-GAP-0094 / KIT-GAP-0095 logged:** step-6 version bump doesn't seed the kit's own RELEASES.md per-patch heading or all ARD version anchors (Technical\_Overview changelog range, user-guide title), so PSK001 + the step-8 critic flag them until manually written.
- (v0.6.75–v0.6.88 per-patch history: see `CHANGELOG.md`.)

## v0.6.90 — Session-health monitor (context-drift badge + /clear reminder)

- **KIT-GAP-0090 — context-window health monitor:** `psk-session-monitor.sh` Stop-hook measures live context occupancy from the Claude Code transcript usage (`input_tokens` + `cache_creation_input_tokens` + `cache_read_input_tokens`, jq-free + schema-verified) → 3-level `ctx:` breadcrumb badge (🟢 <50% · 🟡 50-79% · 🔴 ≥80%) rendered next to `opt:`, countering the lost-in-the-middle attention decay that dilutes a session's early rules.
- **Passive vs active:** the badge renders continuously; a de-duped `/clear` recommendation banner fires in the red zone only (yellow is badge-only — no nag; re-arms after a clear).
- §Session Health Indicator rule added to `portable-spec-kit.md`; wired as a session-stop hook by `psk-install-hooks.sh`; fail-safe (any error or missing data exits 0 silently, never blocking a turn).

## v0.6.89 — Kit-wide model-policy (cost/perf model selection)

- **KIT-GAP-0089 — role→model policy:** `.portable-spec-kit/model-policy.yml` maps each agent role to a model + `psk-model-policy.sh` resolver (phase→role→model, with a built-in fallback) + `psk-spawn.sh` injecting and loudly surfacing `MODEL=` at every spawn (attention-decay-proof — re-printed at the unskippable spawn protocol, not memory-driven) + a PSK046 wiring-integrity gate + manifest registration.

- **Roles:** qa/critic → sonnet (high-fan-out, bounded); dev/qa\_synthesis/implement → opus (correctness-critical); default/plan → inherit.

### v0.6.74 — Sub-Agent Prompt Fidelity 9th Reliability Layer

- Closes KIT-GAP-0055 — sub-agent semantic-deviation class. The kit's 8 prior layers govern spawn routing, workflow ordering, plan schema, command form — but not what's inside a sub-agent's prompt. Layer 9 closes this structurally at every spawn depth.
- **kit-rules.yml manifest:** 19 actionable rules covering verdict authority, scope / severity / bucket / disposition classification, spawn / workflow rules, citation rules, convergence criteria.
- **psk-rule.sh lookup helper:** programmatic verbatim rule access. Sub-agent prompts use script output, not paraphrase.
- **psk-prompt-lint.sh + PSK042 sync-check:** validates `kit_rule_citations:` frontmatter + body verbatim quotes. Advisory in `--quick`, escalates in `--full`.
- **Mandatory preamble + decision-write ban:** "Before any decision, run `psk-rule.sh` lookup. If prompt body contradicts kit verbatim, kit wins. You declare evidence; kit declares decisions." Sub-agent verdict field ignored; verdict computed from gates + findings.
- **psk-spawn.sh prompt-validation gate:** lints prompt via `psk-prompt-lint.sh --strict` BEFORE marking `AWAITING_SUBAGENT`. Lint failure → spawn REFUSED (exit 4), no inline-fallback. Recursive at every spawn depth.
- **Tests:** 23/23 in Section 95. 97/97 prompts lint clean after migration retrofit.

### v0.6.73 — Convergence-prep + citable\_quote Mandate

- KIT-GAP-0054: dim-agent `citable_quote` MUST include `file:line` prefix. Without it gate-13 audit-completeness flags multi-author dispatch passes as synthesis-suspect (signature: `citable_quote_coverage <70%` ).
- 4 stale test fixture fixes: f83 AC2 SAMPLE\_PASS auto-select most recent pass dir, 84.9 bootstrap threshold 50→30, 87.12 open-findings tolerance 0→≤10, G3 grep accepts both legacy and current relaxation-message wordings.
- Tests: 2764/2764 passing.

### v0.6.72 — QA-Orchestrator Multi-Author Dispatch

- Closes 6-pass MAJOR QA-ORCH-COVERAGE finding via §Spawn Fidelity Standard Spawn Recipe.
- KIT-GAP-0052 P1: `psk-spawn.sh request-multi` + `complete-multi` primitives. One phase, N parallel sub-agent spawns via YAML manifest. State machine marks `AWAITING:MULTI_SUBAGENT_SPAWN` .

- KIT-GAP-0052 P2-3: QA-Orchestrator three-path dispatch (multi-author > Task-inline > single-author). New `qa_agent.spawn_mode` config field. `spawn-qa.sh` awk parse + `qa-task.md` plumbing.
- Tests: 20/20 in Section 94. Verified end-to-end on cycle-26-pass-002 (3 parallel dim-agents → GRANTED).

### **v0.6.71 — PSK005 Recursion Guard**

- KIT-GAP-0051: 3-way infinite recursion between `psk-sync-check` ↔ `test-release-check` ↔ `97-kit-fidelity`. Observed 2.6K+ procs, 8h `gates.sh` hang.
- Fix: `test-release-check.sh` exports `PSK_IN_TEST_RELEASE_CHECK=1` sentinel; `check_rft_gate()` short-circuits when sentinel set.
- Verified: test-release-check 8h hung → 24s clean, sync-check 38/38, autoloop unblocked.

### **v0.6.70 — Sync-Check Perf Sweep + Gates Trap**

- Sync-check wall-clock 60-90s → 23s clean. 6 KIT-GAPs closed.
- KIT-GAP-0050: provisional `verdict.md` pre-write before `gates.sh` runs. L2 EXIT trap sees verdict + skips INTERRUPTED stamp.
- KIT-GAP-0048: PSK031 CYC-suffix peel-order two-phase strip.
- KIT-GAP-0049 #1-4: template-quality mtime cache, PSK040 awk single-pass, gate-3 dropped --with-tests, dev-self-verify timeout wrap.

### **v0.6.69 — Autoloop Convergence Root-Cause Fixes**

- KIT-GAP-0017/0019: closes the cycle-22-pass-004 reattachment chase that consumed v0.6.67/v0.6.68 sessions. 6 root-cause bugs fixed in autoloop state machine.
- PSK031 awk-single-pass refactor — eliminated ~2,250 cumulative subshells.
- KIT-GAP-0010: widened preconditions HEAD-pattern walk-back accepting intermediate chore commits.

### **v0.6.68 — Kit-Machinery Sweep**

- KIT-GAP-0012: PSK031 alias-map canonical\_id lookup before flagging duplicate fingerprints.
- 5 KIT-GAPs closed (0009 through 0014).

### **v0.6.67 — Recursion Fixes: Legitimate-Exception Taxonomy**

- **Second production validation of §Kit Fidelity friction-as-feedback loop:** Cycle-22 reflex autoloop surfaced 4 recursive false-positives (PSK040, PSK041, Dim 27) + the KIT-GAP-0013 hook-installer mangling bug. v0.6.67 closes them via a

shared **legitimate-exception taxonomy** so kit rules can distinguish operator-controlled non-pending states from real workaround anti-patterns.

- **PSK041 disposition enum extended** from `{pending, kit-fixed, escalated}` to `{pending, deferred, bypassed, outside-repo, escalated, kit-fixed}`. Operator-controlled non-pending states skip enforcement; only `pending` markers fire.
- `psk-kit-cmd.sh --log-gap` **new flags:** `--defer <target>`, `--bypassed <flag>`, `--outside-repo <path>`. Each sets disposition + `defer_target` in the pending marker so PSK041 sees the operator's explicit intent.
- **PSK040 commit-shape whitelist:** Legitimate release-ceremony and KIT-GAP-closer commit subjects (`release: vX.Y.Z`, `fix(KIT-GAP-NNNN):`, `vN.M.P:`, `chore(release):`, `docs(release):`) skip marker-pattern detection.
- **Dim 27 documented-fallback bypass:** The `single_author_write_timestamp` signature skips when `qa-usage.yaml` declares `mode: single-author-fallback-documented`. The orchestrator's inline write is a documented operator-opted fallback, not a synthesis attack.
- `psk-install-hooks.sh` **hardening (closes KIT-GAP-0011 + KIT-GAP-0013):** New `safe_install_hook()` helper with `bash -n` syntax check before atomic `tmp→mv`, restore from `.psk-backup` on failure, body validation before splicing. Prevents the dangling `for candidate in \` mangling that broke workspace pre-commit in cycle-22-pass-004.
- **6 flow docs synced to v0.6.66 counts:** 38 PSK rules (PSK001-PSK041), 13 mechanical gates, 28 reflex dimensions, 2720 framework tests, 85 features.
- **Defers:** KIT-GAP-0012 → v0.6.68 (PSK031 dedup-vs-alias gap — does not consult findings-registry `canonical_id` before flagging duplicate fingerprints).
- **Tests:** 2865 (2720 framework + 145 benchmarking), 0 failures.

## **v0.6.66 — Kit-Machinery Hardening (5 KIT-GAPs closed)**

- **First production validation of §Kit Fidelity friction-as-feedback loop:** v0.6.64 shipped the §Kit Fidelity layer; v0.6.65 exercised it across two release ceremonies + an autoloop debug session, surfacing 7 distinct kit-machinery bugs. v0.6.66 closes 5 of them at the kit-machinery level — no workarounds, all referenced by KIT-GAP-NNNN id.
- **New helpers:** `psk-version-bump.sh` (patch-bump with per-run idempotency marker — closes KIT-GAP-0002, KIT-GAP-0007), `--next-cycle` escape hatch in `reflex/run.sh` for autoloop recovery (closes KIT-GAP-0006), PSK041 sync-check rule auditing pending-disposition markers (closes KIT-GAP-0008).
- **Tests:** 2865 (2720 framework + 145 benchmarking), 0 failures.



## v0.6.64 — Kit Fidelity (8th Reliability Layer)

- **§Kit Fidelity — 8th reliability layer:** New MANDATORY section in `portable-spec-kit.md`. Every kit-command invocation routes through `agent/scripts/psk-kit-cmd.sh` wrapper with no inline-fallback branch (same asymmetry as `psk-spawn.sh` from v0.6.60). Two principles: (1) canonical default form — non-canonical variants require `--rationale "<text>"` ( $\geq 20$  chars, user-authored) committed to `agent/.kit-deviation-log`; (2) friction = kit bug — workarounds are forbidden; agent files `KIT-GAP-NNNN` in `agent/.kit-gap-log` instead.
- **New wrapper** `agent/scripts/psk-kit-cmd.sh` : routes every canonical kit command via a single entry point. Subcommands: `reflex | prepare-release | refresh-release | init | orchestrate | feature-complete | new-setup | existing-setup | run-plan`. Utilities: `--list` (inventory display), `--check` (dry-run classification), `--log-gap` (file friction without running). Pauses with `AWAITING_RATIONALE` on non-canonical invocation.
- **New canonical-command inventory** `.portable-spec-kit/kit-commands.yml` : data-driven, `schema_version: 1`. Lists each command's canonical default + non-canonical variants with `why_non_canonical` reasons + `marker_commit_patterns` for PSK040 audit. Operators extend by editing YAML, not the wrapper.
- **Two committed audit-trail files:** `agent/.kit-deviation-log` (append-only TSV per `--rationale` invocation: *ts CMD FLAG hash text*) and `agent/.kit-gap-log` (auto-incremented `KIT-GAP-NNNN` friction ledger: *ts ID CMD FRICTION FIX*).
- **New sync-check rule PSK040:** audits git log since the §Kit Fidelity introduction commit for marker-shaped commits without matching deviation-log entries. ADVISORY in `--quick` mode (PostToolUse hook), ERROR in `--full` mode (PreCommit hook). Defense-in-depth: wrapper PREVENTS at invocation time, PSK040 DETECTS post-commit if the wrapper was bypassed.
- **New flow doc** `docs/work-flows/30-kit-fidelity.md` + **JIT-loaded skill** `.portable-spec-kit/skills/kit-fidelity.md` + **test section** `tests/sections/97-kit-fidelity.sh` (33 assertions across 13 test groups, wired into orchestrator) + **ADR-089** in `agent/PLANS.md` + **F85** in `SPECS.md`.
- **§Reliability Architecture overview text:** bumped from "six enforcement layers" to "eight enforcement layers" (the 7th — §Workflow Declaration — was already shipped in v0.6.62 but the overview text wasn't synced).
- **Why this layer:** recurring agent failure 2026-05-30/31 where the agent picked `reflex/run.sh single` instead of canonical autoloop, then empty marker commit instead of `prepare release`, then stopped after DENIED single-pass instead of iterating. Each deviation had operational excuse; "follow kit designs" rule existed but got rationalized past under pressure. Same class as the synthesis-substitution failure that drove §Spawn Fidelity in v0.6.60.

- **Emergency bypass:** `PSK_KIT_FIDELITY_DISABLED=1` logged to `agent/.bypass-log` per PSK027.
- **Tests:** 2720 framework + 145 benchmarking = 2865. 0 failures. 85 features, 100% R→F→T coverage on v0.6.64 additions. 30 flow documents (added 30-kit-fidelity.md).

## v0.6.63 — Reflex Convergence + Workflow Declaration Flow Doc

- **QA-C22-04 fix (CRITICAL):** `reflex/lib/findings-registry.sh` honors `REFLEX_FINDINGS_REGISTRY` + `REFLEX_HISTORY_DIR` env overrides so registry tests no longer mutate the committed registry. Section 84 fixtures isolate via temp registry — suite is now deterministic.
- **QA-C22-07 fix (CRITICAL):** `reflex/lib/update-eval-trace.sh` emits the "Cycle-numbering note" callout in the generated `REFLEX_EVAL_TRACE.md` header so per-pass regeneration no longer wipes it. Test 82.19 deterministic.
- **QA-C22-09 fix (MINOR):** PSK032 grandfather exemption for cycles carrying `migration-note.md` (documented historical mis-numbering). Regression test 82.21.
- **QA-C22-08 fix (MINOR):** `reflex/lib/doc-code-diff.sh` honors `# doc-coverage-exempt:` header marker on 4 internal mechanical helpers — precise (only marked scripts opt out, no blanket silencing). 2 regression tests.
- **Added** `docs/work-flows/29-workflow-declaration-schema.md` : documents the 7th reliability layer introduced in v0.6.62 (PSK034 + PSK035 + Class A/B/B'/B-plan-driver/C taxonomy + phases.yml schema + dispatch contract).
- **Changed:** stale `reinit` references in `13-release-workflow.md` repointed — 5 executable workflows (release, feature-complete, init, new-setup, existing-setup) with `reinit` documented as folded into idempotent `init`.
- **Tests:** 2687 framework + 145 benchmarking = 2832. 0 failures. 3 new regression tests (82.21 PSK032 grandfather + 2 doc-coverage-exempt in `04-reflex.sh`). 30 flow documents (29 user-facing + 1 template).

## v0.6.62 — Command-Model Redesign + Dispatcher Migration

- **Two-command model:** idempotent `install` + `init`. `init` becomes the registry-driven CREATE-or-REFRESH conformance engine, content-loss-protected. The retired `reinit` command is folded into `init` — one command covers empty-project create AND existing-project refresh.
- **Unified** `orchestrate build` : one workflow for new + existing projects. Each phase is idempotent create-or-update so the same dispatch handles both modes. `--update` / `--retrofit` modes removed.
- **PSK034 + PSK035 sync-check rules implemented** (were documented in v0.6.57 but never coded): PSK034 audits workflow-router↔phases.yml mapping; PSK035 validates the `phases.yml` schema (schema\_version, workflow, phases array per-phase fields, dependency-graph validity).

- `psk-run-plan.sh` **delegates start/next/resume/retry to** `psk-dispatch.sh --plan` : every plan-phase spawn now routes through `psk-spawn.sh` — extends §Spawn Fidelity coverage to plan execution. Compat-mode preserved for legacy narrative plans (one-shot then convert).
- **Reflex resolved as monolithic-by-design (Class B')**: `reflex/run.sh` 's iterate-until-convergence control flow (GRANTED/DENIED branching + stateful L1-L6) does not fit the dispatcher's linear `phases.yml` model. Documented as architecturally-distinct, NOT pending migration. The §Spawn Fidelity protocol already covers reflex's sub-agent spawns via `spawn-qa.sh` + `spawn-dev.sh` .
- **Class A / B / B' / B-plan-driver / C workflow taxonomy** formalized so retrofit scope is unambiguous.
- **Tests:** 2687 framework + 145 benchmarking = 2832. 0 failures.

### v0.6.61 — Strict-GRANT Convergence

- **Strict-GRANT convergence discipline:** reflex autoloop no longer advances cycles on findings-zero-but-DENIED states — only on explicit GRANTED verdicts. Closes the false-convergence class where a pass with 0 NEW findings still left prior-cycle findings unresolved.
- **Verdict integrity hardening:** reflex/lib write-verdict structural field emitting explicit `verdict: GRANTED|DENIED` . Per-pass `.iter-status.yml` audit trail tracks PENDING → RUNNING → COMPLETE-GRANTED / COMPLETE-DENIED transitions.
- **Convergence-audit gate (9th mechanical gate):** `reflex/lib/gates.sh` fails any pass that left an `INTERRUPTED` verdict without an `operator-recovered` annotation. Recovery requires explicit `--recover-from-abort <pass-id>` flag.
- **Tests:** 2687 framework + 145 benchmarking = 2832. 0 failures.

### v0.6.60 — Spawn Fidelity Hardening

- **§Spawn Fidelity — 6th reliability layer:** New MANDATORY section in `portable-spec-kit.md` . Every sub-agent invocation routes through `agent/scripts/psk-spawn.sh` with no inline-fallback branch. P11 — Spawn Fidelity principle added to `agent/PHILOSOPHY.md` .
- **HF1 — Reflex spawn-qa/spawn-dev retrofit:** `reflex/lib/spawn-qa.sh` and `reflex/lib/spawn-dev.sh` route through `psk-spawn.sh` , inheriting the no-inline-fallback contract automatically.
- **HF1b — Orchestrate --update per-phase split:** `psk-orchestrate.sh --update` U3-U10 split from single batch into per-U-phase spawns. Closes the 15h-batch context-overflow risk surfaced in v5 Phase C.
- **HF2 — psk-critic-spawn retrofit:** All 6 critic templates (STEP\_9\_VALIDATION, FEATURE\_COMPLETE, INIT, REINIT, NEW\_SETUP, EXISTING\_SETUP) route through the spawn-fidelity wrapper.

- **HF3 — Persistent retry queue:** `agent/.workflow-state/retry-queue.yml` + `psk-retry-queue.sh` CLI. Exponential backoff (5min → 15min → 45min → 2h → 6h → AWAITING\_HUMAN\_ARBITRATION). Survives session ends.
- **HF4 — Resume-on-session-start MANDATORY rule:** First agent action on every session entry runs `psk-resume-bootstrap.sh`. PSK029 sync-check detects sessions that started without bootstrap.
- **HF4b — Workflow watchdog + phase idempotency:** `psk-workflow-watchdog.sh` (15min WARN / 1h HUNG / 24h STALE) + phase idempotency contract (gate-check + artifact-check + atomic writes + idempotent commits).
- **HF5/HF6 — Synthesis-detection probe + 13th gate:** `reflex/lib/check-audit-completeness.sh` 6-signature analysis. Gate 13 `audit-completeness` blocks at synthesis-confirmed verdict. PSK026 sync-check mirrors for critic-result.md files.
- **HF7/HF7b — Dim 27 + Dim 28:** Recursive synthesis audit of previous pass's findings.yaml (Dim 27); grep-detects kit scripts doing inline AI work without `psk-spawn.sh` (Dim 28, 4 patterns). Both surface as scope:kit findings routed through PKFL.
- **HF8/HF8b — §Spawn Fidelity + Standard Spawn Recipe skill:** New skill at `.portable-spec-kit/skills/spawn-fidelity.md` codifies the recipe: 7-section prompt template, 4-field artifact template, workload-driven spawn-count mandate, 8-step Dev-Agent fix protocol for Dim 28 findings.
- **HF9 — PSK027 bypass-tamper-detection:** `psk-bypass-log.sh` logs every `PSK_*_DISABLED=1` invocation to `agent/.bypass-log`. PSK027 fires WARNING at 1-2 bypasses / ERROR at 3+ in 24h.
- **HF0 — Cascade-as-user-update purge + PSK028:** Cascade-as-user-update path purged from 5 kit sites. Canonical user-project update is now `bash agent/scripts/psk-orchestrate.sh --update`. PSK028 anti-pattern rule flags reintroduction.
- **SPECS F77-F84 added:** One feature per HF mechanism. 28 flow documents (added `28-spawn-fidelity.md`). New flow doc 28 (Spawn Fidelity).
- **Tests:** 2377 framework + 145 benchmarking = 2522 total. 0 failures. 84 features, 100% R→F→T coverage. 30/30 sync-check rules (PSK026 + PSK027 + PSK028 + PSK029 added).

## v0.6.57 — Plan Execution Protocol + UI Completeness Gate

- **New MANDATORY §Plan Execution Protocol in** `portable-spec-kit.md` : 5th reliability layer (plan-shape gate). Every plan executed via sub-agent spawning conforms to `schema_version: 1` + `phases:` frontmatter. Legacy plans run once under `compat_mode: true`, then MUST convert before next `start`.
- **Three new templates in** `.portable-spec-kit/templates/` : `plan-executable.md`, `plan-prompt.md`, `plan-artifact.md` — all pass the 7-criterion Template Quality Bar.

- **New script** `agent/scripts/psk-run-plan.sh` **(35th kit script)**: schema-aware driver. Commands: `start` / `next` / `status` / `resume` / `retry` / `--convert` / `--validate` / `--health` / `abort` . Refuses non-conformant plans (one-shot compat-mode exception). No inline-fallback branch — structurally only spawn/retry/abort paths forward.
- `psk-plan-save.sh` **body-preservation fix + schema validation on** `approve` : today's regression erased 233 lines of plan body — fixed via tempfile + frontmatter-only rewrite. `approve` runs schema validation (PSK024 codes N/P/F/D/L); new `--validate-schema` subcommand.
- **PSK024 sync-check rule** (27th rule, 11 sub-codes V/P/I/N/R/A/G/C/D/M/X) lints every `agent/plans/*.md` . PostToolUse warn + pre-commit error.
- **§Workflow Fidelity B2 retrofit**: `register-gate` / `verify-gate` API on `psk-workflow-state.sh` + 9 workflows wired. `mark-done` refuses unverified phases.
- **New script** `agent/scripts/psk-ui-completeness.sh` **(36th kit script)**: stack-aware 10-category UI audit. Sub-codes — P primitives, L layout, D dark-mode, S loading/empty/error states, A a11y, T design-tokens, F per-feature pages, I input-feedback, R responsive, E empty-shell anti-skeleton. Commands: `default` / `--check` / `--strict` / `--json` / `--feature` . Bash 3.2 compatible.
- **PSK025 sync-check rule** (28th rule) — wraps `psk-ui-completeness.sh --json` . Bypass: `PSK_UI_COMPLETENESS_DISABLED=1` .
- **Orchestrator** `--update` **U6.5** — new UI-completeness backfill phase between sync-check-config and features. Sub-agent reads `--json` output, backfills every PSK025 violation with full implementations, re-verifies. No skip path.
- **Reflex Dim 26 + 12th mechanical gate** `workflow-fidelity-completeness` : audits workflow-state ledger + UI completeness + per-phase gate discipline + spawn-protocol fidelity. Block at MAJOR (configurable).
- **New flow docs 26 + 27**: `26-plan-execution-protocol.md` + `27-ui-completeness-gate.md` .
- **New skill** `.portable-spec-kit/skills/plan-execution.md` — agent-facing operational guide.
- **Motivating incidents closed**: (1) plan body erasure regression — catalyst for stricter schema enforcement. (2) v5 skeletal UI — empty-shell pages now auto-flagged as PSK025-E.
- **Tests**: 2180 framework + 145 benchmarking = 2325 total. 390 reliability tests across Sections 61-67. 0 failures.

## **v0.6.56 — Workflow Fidelity: the agent executes the kit, never overrides it**

- **New MANDATORY §Workflow Fidelity in** `portable-spec-kit.md` : the 4th reliability layer — behavioral, structurally enforced. When the kit defines a process

(workflow, phase, gate, rule, principle, test) the agent executes it faithfully and completely. Rate/context limits → pause-and-resume, never reduce-scope.

- **New principle** P10 – Workflow Fidelity in `agent/PHILOSOPHY.md` — evidence base: searchsocialtruth-v5 shipped a skeletal UI behind a "20/20 features, 110 tests" report.
- **New script** `agent/scripts/psk-workflow-state.sh` (**33rd kit script**) + `agent/scripts/psk-spawn.sh` (**34th kit script**) with no inline-fallback branch.
- **9 executable workflows retrofitted**; new flow doc 25; 18 regression tests in section 61.

## v0.6.55 — Source Layout Policy + 10 new templates + flow doc 24

- **New §Source Layout Policy** in `portable-spec-kit.md` : Stack-driven two-step decision — (1) Template choice (Template 1 colocated, 3 separate services, 4 monorepo, 5 library, 6 CLI) via PSK022a; (2) `src/` opt-in subdirs ( `core/` + `shared/` always; `ui/` , `api/` , `integrations/` , `platform/` per Stack table) via PSK022b. No empty mandatory dirs.
- **10 new templates** in `.portable-spec-kit/templates/` (total 23, all pass 7-criterion Quality Bar): reflex-prompt, test-section, skill, ard-html, kit-finding, research, adr, handoff, src-layout, src-subdir-readme.
- **New script** `psk-scaffold-src.sh` (**32nd kit script**): reads PLANS.md Stack table → derives opt-in subdirs → creates with READMEs. Idempotent. `--check` dry-run.
- **New flow doc 24** — `docs/work-flows/24-template-driven-creation.md` : documents template-driven creation lifecycle + PSK022a/022b/023 gates.
- **Tests**: 2207 framework + 145 benchmarking = 2352 total.

## v0.6.54 — Template Quality Bar + externalized templates dir

- **New §Template Quality Bar** in `portable-spec-kit.md` : 7 mandatory criteria every kit-shipped template must pass — stack-agnostic, domain-agnostic, scale-agnostic, useful-and-complete (refined wording: useful > generic, no vacuous "fill in X" placeholders), lifecycle-aware, self-documenting ( `<!-- TEMPLATE-KIND / GENERICITY / LAST-AUDITED -->` header), round-trippable (scaffold passes sync-check zero-edit).
- **Externalized templates dir** `.portable-spec-kit/templates/` : 13 templates extracted from inline `.portable-spec-kit/skills/templates.md` (633 lines): 11 agent file templates (REQS, SPECS, PLANS, DESIGN, RESEARCH, TASKS, RELEASES, AGENT, AGENT\_CONTEXT) + design-plan + flow-doc + plan + project-readme.
- **New script** `agent/scripts/psk-template-quality.sh` (**31st kit script**): 7-criterion lint with JSON output ( `--json` ), bulk audit ( `--all` ), strict gate mode ( `--all --strict` ), audit log writer to `.audit-log.yaml` .
- **PSK023 sync-check rule**: `check_template_quality()` runs `psk-template-quality.sh --all --strict` on every `psk-sync-check.sh --full` . Blocks templates that fail the bar at the mechanical gate.



- **Predicted failure caught & fixed:** `agent-AGENT-template.md` initially failed criterion 1 (stack-specific examples in normative sections — Hono / FastAPI / Zod / Pydantic) and 7 ( `TBD` placeholders without REQUIRED marker). Fixed by extracting stack-specifics to `### Examples (stack: typescript-node)` + `### Examples (stack: python)` blocks and marking placeholders `<!-- REQUIRED — replace before commit -->`.
- **Phase A0 of the Template Completeness plan:** v0.6.54 established the 7-criterion quality bar against the existing templates; v0.6.55 added 10 new templates + the Source Layout Policy on top of that proven bar.
- **README install count:** 30 → 31 reliability scripts + new "13 file-class templates" callout.
- **Tests:** 2206 framework + 145 benchmarking = 2351 (unchanged — additive feature).

### v0.6.53 — Plan-Save Protocol: interruption-resilient plan persistence

- **New rule §Task Tracking / Plan-Save Protocol:** any conversation-drafted plan ( $\geq 3$  markdown headings OR  $\geq 30$  lines AND a plan-shaped prompt) MUST persist to `agent/plans/YYYY-MM-DD-<slug>.md` as the agent's first tool call of the same turn, before any implementation tool calls. Lifecycle states live in frontmatter — `draft` → `approved` → `executing` → `done` | `abandoned`.
- **New script** `agent/scripts/psk-plan-save.sh` (**30th kit script**): subcommands `save` / `approve` / `start` / `done` / `abandon` / `list` / `show`. Idempotent — re-saving a plan with the same slug rewrites the body and refreshes the `updated:` field while preserving lifecycle status.
- **Motivating incident:** the 705-line `kit-v0639-and-sst-v2.md` plan lived only in `~/.claude/plans/` (ephemeral IDE storage) and was nearly lost on context compact. Plan-Save Protocol closes this gap.

### v0.6.51 — Reflex cycle-10/pass-002 merge: doc-code-diff fixes + F16 test hardening + K3/K7 cross-refs

- **Reflex cycle-10/pass-002 merged:** 119 files committed — all pass-002 artifacts (findings.yaml, signoff.md, verdict.md, dev-trace.md, 72 F-test plans, coverage.md, gates-result.md, .cycle-meta).
- **doc-code-diff.sh template exclusions:** Added false-positive filter patterns for placeholder file names used in kit documentation examples: `foo.sh`, `psk-foo.sh`, `script-name.sh`, `test_f*_feature_name.*`, `f*-feature-name.*`. Eliminates spurious MAJOR noise from instructional literals.
- **F16 test hardened (exact count):** `tests/features/f16-work-flow-docs.sh` assertion changed from `>=` to `==`; `00-template.md` excluded from numbered workflow doc count — test now asserts exactly 23 numbered docs.



- **23-project-update-workflow.md K3/K7 additions:** Added K3 (design-plan and stub creation for U3/U4), enhanced K6 cross-references, added K7 (HANDOFF.md structure for U10).
- **Tests:** 2206 framework + 145 benchmarking = 2351 (+1 from doc-code-diff basename fix).

## **v0.6.50 — Workflow template standard: PSK019 gate + all 23 flow docs standardised + Gate 1 scope fix**

- **Workflow template standard (PSK019):** New `docs/work-flows/00-template.md` establishes the canonical 7-section template (Overview, Flow Diagram, Steps, Key Rules, Edge Cases, Related Flows, Bypass) with 63-char ASCII box standard. All 23 existing flow docs updated to comply — every doc now has `## Overview`, `## Flow Diagram`, and `## Key Rules` sections.
- **PSK019 gate in psk-sync-check.sh:** `check_flow_doc_template()` enforces the 3 required sections on every `docs/work-flows/*.md` except `00-template.md` itself; PSK015 and PSK003 exclude it from doc-count and file-count checks.
- **Box line alignment (63-char standard):** All ASCII box lines in the 23 flow diagram boxes aligned to exactly 63 display characters — 80 misaligned lines fixed across 13 files; enforced by infrastructure test.
- **Reflex Gate 1 scope fix:** `reflex/lib/preconditions.sh` Gate 1 changed to `git status --porcelain -- .` — scopes clean-tree check to `PROJ_ROOT` subtree, preventing false positives in workspace-level git repos.
- **Tests:** 2191 framework + 145 benchmarking = 2336 total.

## **v0.6.49 — resolve\_kit\_root bug fix: path-spaces stripping + heuristic false-positive**

- **K1 path-spaces bug fixed:** `resolve_kit_root()` used `tr -d '[:space:]'` stripping embedded spaces in directory names (e.g. "Aqib Mumtaz"). Fixed to `sed 's/[[[:space:]]*$//'` which strips only trailing whitespace.
- **Heuristic false-positive fixed:** Any user project with a real copy of `portable-spec-kit.md` was being identified as the kit root. Fixed by adding `[ -f "$kit_candidate/install.sh" ]` check — only a real kit directory has `install.sh`.
- **doc-23 updated:** `23-project-update-workflow.md` resolution order item 3 now documents the `install.sh` check requirement.
- **Tests:** 2191 framework + 145 benchmarking = 2336 total.

## **v0.6.48 — Orchestrator --update mode + Phase 9 mandatory reflex + KIT\_ROOT resolution + polyglot stack detection + R→F→T stubs + project-update flow doc**

- **K1 — KIT\_ROOT persistence:** `install.sh` writes `kit_source_path`, `kit_version`, and `kit_installed_at` to `.portable-spec-kit/config.md` so any project can resolve its kit source without user intervention.
- **K2 — Phase 9 mandatory reflex:** `psk-orchestrate.sh` Phase 9 (reflex-audit) now has NO skip path. Reflex is auto-installed if missing via `resolve_kit_root()`; convergence loop runs until GRANTED.
- **K3 — --update mode (U0–U10):** `psk-orchestrate.sh --update` brings any speckit project to current kit in a single command: `script sync → mandate audit → reflex install → design plans → feature stubs → ARD/flows → sync-check config → features → release-prep → reflex GRANTED → handoff`. Replaces the retired manual cascade pattern.
- **K4 — R→F→T stubs in orchestration:** Phase 7 sub-agent prompt updated to require `tests/features/` anchor stubs alongside runnable tests for full R→F→T traceability.
- **K5 — Polyglot stack detection:** `mandate-audit.sh` now detects `backend/` + `frontend/` projects and resolves them as `polyglot-python-node` instead of `unknown` — mandates no longer silently skipped.
- **K6 — summary.csv GRANTED path:** `loop.sh` GRANTED terminator now calls `score.sh` to write the summary row; cycle-09 row backfilled.
- **K7 — Flow doc 23:** `docs/work-flows/23-project-update-workflow.md` documents the --update mode U0–U10 phases and formally retires the manual cascade pattern.
- **Tests:** 2191 framework + 145 benchmarking = 2336 total.

## **v0.6.44 — Kit-evolution phases 1–6: config.md, template enrichment, F71/F72, test infra fixes**

- **`.portable-spec-kit/config.md`** created and wired into `install.sh` (auto-created on install), `psk-sync-check.sh` (PSK015-CONFIG gate), `mandate-audit.sh` (MANDATE-CONFIG-MISSING MAJOR probe), and `tests/sections/05-mandate-compliance.sh` (T4/T5 regression tests).
- **Template enrichment ( `skills/templates.md` ):** All 9 agent file templates updated with gold-standard sections — REQS (approval tracking + requirement types + source columns), SPECS (scope-change record + tests/features stub path), AGENT.md (Security + Definition of Done), TASKS.md (QA Findings + Evidence field), RESEARCH.md (Open Research Questions), AGENT\_CONTEXT.md (tests/features/ + config.md in File Structure).
- **F71 — Framework Internal Consistency:** 6 regression tests in `tests/features/f71-framework-consistency.sh`.

- **F72 — Template Structure Completeness:** 15 regression tests in `tests/features/f72-template-structure.sh` .
- **PSK016 gate + Dim 25 extension + test infra fix. Tests:** 2336 total.

**v0.6.43 — Flow doc consistency sweep: step counts, gate count, dimension count, test count**

- **Step count fix** ( `11-spec-persistent-development.md` ): "8-step gate" → "10-step gate" corrected in two places — box diagram and Pipeline Sync Rules table.
- **Gate count fix** ( `17-reflex.md` ): "10 gates" → "11 gates" — PKFL added G1 (kit-evolution-file-scope) + G2 (kit-genericity-proof) in v0.6.41.
- **Dimension count fix** ( `18-project-orchestration.md` ): "24 dimensions" → "25 dimensions" — Dim 25 Mandate-Compliance added in Loop-Iter-2.
- **Test count fix** ( `19-kit-evolution-gauntlet.md` ): "1836 tests" → "1982 tests" (1837 framework + 145 benchmarking).
- **Test infrastructure:** flow count check in `tests/sections/01-infrastructure.sh` updated 21 → 22 to match README set in v0.6.42.
- **Tests:** 1837 framework + 145 benchmarking = 1982 (no count change).

**v0.6.42 — README consistency sweep: script count + flow table row 22**

- **PSK013 fix:** README Quick Start "Installs:" updated from 27 → 29 scripts — `psk-generate-ci.sh` and `psk-generate-user-guide.sh` added in v0.6.41 were not reflected.
- **PSK015 fix:** README Flows table row 22 added for `22-project-kit-feedback-loop.md` ; Documentation section reference updated from "21" → "22 step-by-step flow diagrams". ARD Key Highlights updated: flow docs 21→22, scripts 27→29.
- **Tests:** 1837 framework + 145 benchmarking = 1982 (no count change).

**What the Framework File Contains**

The `portable-spec-kit.md` file is the single source of truth for your development methodology. It contains:

| Section            | What It Governs                                                                           |
|--------------------|-------------------------------------------------------------------------------------------|
| User Profile       | Personalized AI — GitHub auto-detect, communication style, working pattern, AI delegation |
| Git & GitHub Rules | When to commit, push, ask before critical ops                                             |

|                              |                                                                                                                                  |
|------------------------------|----------------------------------------------------------------------------------------------------------------------------------|
| Security Rules               | Secret handling, .env practices, code security                                                                                   |
| Versioning                   | Two-level: framework v0.N.x (each publish) + release v0.N (milestones), auto-restructure on pull                                 |
| Task Tracking                | Add tasks FIRST, then work. Module-based organization.                                                                           |
| Testing Standards            | Coverage targets, edge case checklist, mock rules, self-validation                                                               |
| Code Quality                 | Review checklist, naming conventions, deployment checklist                                                                       |
| Error Handling               | Structured errors, logging, error boundaries, user-friendly messages                                                             |
| Branch & PR Workflow         | Feature branches, PR format, squash merge, CI-before-merge requirement                                                           |
| CI & Community Contributions | GitHub Actions ci.yml template, stack-aware test commands, branch protection, PR workflow, contribution validation               |
| Spec-Based Test Generation   | Forward-flow stub generation, per-feature acceptance criteria format, stack-aware stubs, stub completion gate                    |
| Dependencies                 | Bundle size checks, lock files, audit, avoid unnecessary deps                                                                    |
| Project Templates            | 6 agent files + README + 8 source code structures (Web, Python, Mobile, Android, iOS, Full Stack, Full Stack + Mobile, Document) |
| Auto-Scan                    | Detects existing projects, creates missing files, restructures to standard — 16 work flow docs                                   |
| Agent Behavior               | Guide don't enforce, silent tracking, retroactive spec-filling, $R \rightarrow F \rightarrow T$ traceability                     |

## Multi-Agent Support

One source file — `portable-spec-kit.md` — works with every AI coding agent via symlinks:

| Agent          | File Created                                 | How                                         |
|----------------|----------------------------------------------|---------------------------------------------|
| Claude Code    | <code>CLAUDE.md</code>                       | Symlink → <code>portable-spec-kit.md</code> |
| GitHub Copilot | <code>.github/copilot-instructions.md</code> | Symlink → <code>portable-spec-kit.md</code> |

|                 |                             |                                |
|-----------------|-----------------------------|--------------------------------|
| <b>Cursor</b>   | <code>.cursorrules</code>   | Symlink → portable-spec-kit.md |
| <b>Windsurf</b> | <code>.windsurfrules</code> | Symlink → portable-spec-kit.md |
| <b>Cline</b>    | <code>.clinerules</code>    | Symlink → portable-spec-kit.md |

**macOS/Linux:** Uses symlinks — edit `portable-spec-kit.md` , all agents read the update instantly.

**Windows:** Uses copies — re-run install command after editing, or let the AI agent keep files in sync.

# Complete Flow

---

- 1 **Install** — one command downloads `portable-spec-kit.md` + creates symlinks for all agents (Claude, Copilot, Cursor, Windsurf, Cline)



- 2 **Open any AI agent — auto-scaffolds** — detects environment, creates `WORKSPACE_CONTEXT.md`, creates `agent/` with 6 files (`AGENT`, `AGENT_CONTEXT`, `SPECS`, `PLANS`, `TASKS`, `RELEASES`), project dirs (`src/`, `tests/`, `docs/`), `README`, `.gitignore`, `.env.example`, `tests/test-release-check.sh`



- 3 **Work however you want** — "build me X" → agent adds to `TASKS.md`, builds, tests. "What should I do next?" → agent walks through spec-persistent flow. Never blocks.



- 4 **Session ends** — agent updates `AGENT_CONTEXT.md` with progress, decisions, what's next



- 5 **Come back later** — days, weeks, months. Agent reads `AGENT_CONTEXT.md`: "Here's where we left off." Continues exactly where you stopped.



- 6 **Edit the framework** — update `portable-spec-kit.md`, all 5 agent symlinks read the change instantly. Standards evolve with you.

## Project Structure

---

```
your-project/
├─ portable-spec-kit.md      ← Framework file (source)
├─ .portable-spec-kit/      ← Kit config (user profiles, per-user)
│   └─ user-profile/
│       └─ user-profile-{username}.md
├─ CLAUDE.md                ← Symlink (Claude Code)
├─ .cursorrules              ← Symlink (Cursor)
├─ WORKSPACE_CONTEXT.md     ← Auto-created (workspace state)
├─ README.md                ← Auto-created (standard structure)
├─
├─ agent/                   ← Auto-created (project management)
│   └─ AGENT.md              ← Project rules, stack, brand
```

```

|   |   | AGENT_CONTEXT.md   ← Living state (updated every session)
|   |   | SPECS.md         ← Requirements & features
|   |   | PLANS.md          ← Architecture, methodology & research
|   |   | TASKS.md          ← Task tracking (module-based)
|   |   | RELEASES.md       ← Version history
|   |
|   | src/                  ← Your code
|   | tests/               ← Your tests
|   | ard/                  ← Architecture docs (HTML + PDF)
|   | ...

```

## The 6 Agent Files

| File             | Purpose                                                                              | When           |
|------------------|--------------------------------------------------------------------------------------|----------------|
| AGENT.md         | Project-specific instructions — stack, rules, brand, AI config                       | Setup          |
| AGENT_CONTEXT.md | Living state — done, next, decisions, blockers, file structure                       | Every session  |
| SPECS.md         | What to build — requirements, features, acceptance criteria, scope                   | Before dev     |
| PLANS.md         | How to build — architecture, data model, phases, methodology, research, decision log | Before dev     |
| TASKS.md         | Module-based task tracking with progress summary table                               | During dev     |
| RELEASES.md      | Version changelog (categorized), test results, deployment log                        | End of version |

## How to Work in Spec-Persistent Development

### Phase 1: Specification

**Goal:** Define WHAT you're building before touching code.

Start every project by writing `agent/SPECS.md`. This is a conversation between you and the AI:

**You Say**

**AI Does**



|                                   |                                                                                        |
|-----------------------------------|----------------------------------------------------------------------------------------|
| "I want to build a resume editor" | Writes requirements, features table, scope boundaries, acceptance criteria in SPECS.md |
| "Add AI-powered bullet rewriting" | Adds feature to specs, estimates priority, updates scope                               |
| "That's out of scope for v1"      | Moves to "Out of scope (future)" section                                               |

**Rule:** No code is written until SPECS.md has requirements and acceptance criteria.

## Phase 2: Planning

**Goal:** Define HOW you're building it — architecture, data model, phases.

AI writes `agent/PLANS.md` based on approved specs:

| Section             | What AI Fills In                                                               |
|---------------------|--------------------------------------------------------------------------------|
| Stack               | Recommends tech (frontend, backend, DB, hosting) with trade-offs. You approve. |
| Architecture        | System design, component relationships, data flow                              |
| Directory Structure | Exact file/folder layout for the chosen stack                                  |
| Data Model          | Tables, types, relationships                                                   |
| API Endpoints       | Method, path, description for each route                                       |
| Build Phases        | Phased delivery plan — foundation first, features later                        |
| Security            | Auth, CORS, input validation, secret handling                                  |

**Rule:** Deployment is deferred to release time. Don't discuss hosting during planning.

## Phase 3: Execution

**Goal:** Build, test, and track every task.

AI creates `agent/TASKS.md` with module-based breakdown:

```
## Module 1: Project Setup
| # | Task | Status |
```

```
|---|-----|:-----:|
| 1.1 | Initialize Next.js + Tailwind | [x] |
| 1.2 | Create component structure | [ ] |

## Module 2: Core Features
| # | Task | Status |
|---|-----|:-----:|
| 2.1 | Build editor canvas | [ ] |
```

Execution rules enforced by the framework:

| Rule               | What It Means                                                                       |
|--------------------|-------------------------------------------------------------------------------------|
| Tasks first        | When you ask for something, AI adds it to TASKS.md BEFORE starting work             |
| Self-validate      | AI runs tests, fixes failures, only presents stable results                         |
| Test everything    | Every feature gets tests — unit, integration, UI interaction                        |
| Edge cases         | Empty inputs, boundaries, XSS, unicode, error paths — all tested                    |
| Coverage targets   | Backend 98%+, frontend logic 98%+, UI components 85%+                               |
| No broken features | User should NEVER discover something doesn't work                                   |
| Context updated    | AGENT_CONTEXT.md updated after significant work, after commits, and before any push |
| Flows updated      | docs/work-flows/ updated if any work flow changed — flows always match code         |

Phase 4: Tracking

**Goal:** Log what was delivered, tested, and deployed.

At end of each version, AI updates agent/RELEASES.md :

```
## v0.2 — AI Features (March 30, 2026)

### Summary
Added JD-tailored resume, PDF generation, 673 tests.

### Changes
- **Frontend:** JD Tailor modal, version manager, skills tags
```

```
- **Backend:** 3 PDF engines, expand API, fetch-url API
- **AI:** Bullet format detection, deduplication, length matching

### Tests
- 673 tests passing, 95% coverage

### Deployment
- Deployed to: aqibmumtaz.github.io
- Date: 2026-03-30
```

## Daily Workflow Summary

- 1 **Start session** — AI reads agent/AGENT.md + AGENT\_CONTEXT.md. Knows exactly where you left off.  

- 2 **You give instructions** — "add dark mode", "fix scoring bug", "build PDF export"  

- 3 **AI adds to TASKS.md** first, then builds, tests, and shows results with coverage  

- 4 **You review** — approve, adjust, or ask for changes. AI fixes and re-tests.  

- 5 **You say "commit"** — AI commits with descriptive message. "push" sends to remote.  

- 6 **Session ends** — AI updates AGENT\_CONTEXT.md. Next session picks up seamlessly.

**The result:** Every project follows the same spec → plan → build → test → track pipeline. Nothing is ad-hoc. Nothing falls through cracks. Context never lost.

# Agent Behavior — Guide, Don't Enforce

The framework is **not rigid**. The AI agent adapts to how you want to work:

**Core principle:** The user can start anywhere — jump into coding, give direct tasks, ask questions, or follow the full spec-persistent flow. All valid. The agent adapts to the user, not the other way around.

| You Do                   | Agent Does                                                         |
|--------------------------|--------------------------------------------------------------------|
| "Build me a login page"  | Adds to TASKS.md → builds it → tests it → shows results            |
| "Fix this bug"           | Adds to TASKS.md → fixes → marks done                              |
| "What should I do next?" | Walks through spec-persistent process: define specs → plan → tasks |
| "What's the status?"     | Shows from TASKS.md + AGENT_CONTEXT.md                             |
| Come back after weeks    | Reads AGENT_CONTEXT.md → summarizes where you left off             |
| Never wrote specs        | Retroactively fills SPECS.md from what's built — no complaints     |

## Silent Tracking

Even if you never mention the process, the agent:

- Adds every task to TASKS.md before starting
- Marks tasks done when complete
- Updates AGENT\_CONTEXT.md after completing work, committing, and before pushing
- Fills gaps in SPECS.md and PLANS.md retroactively
- Surfaces process naturally: "I've added this to TASKS.md", "Based on SPECS.md, we still have..."

**User's time is sacred.** Agent does 90% of the work — writes specs, plans, tasks, tests, docs. User reviews 10%. Never ask the user to write specs or plans — the agent drafts everything, user approves or adjusts.

## File Creation/Update Rule

This single rule governs how the agent handles ALL managed files ( `WORKSPACE_CONTEXT.md` , `README.md` , and all `agent/` files):

| Scenario                        | Action                                                                               |
|---------------------------------|--------------------------------------------------------------------------------------|
| File does not exist             | Create from standard template, fill in known details                                 |
| File exists but wrong structure | Restructure to match template — <b>retain all existing content</b> , never lose data |
| File already matches template   | Leave as-is                                                                          |

**Never lose content.** When restructuring an existing file, every detail, decision, and note is preserved — just reorganized into the standard sections.

## Customization

### User Profile — Personalized AI Experience

On first session, the agent fetches your GitHub identity and asks 3 quick questions. Press Enter to use recommended, or type your own:

| Question                    | Options                                                                                                                                                                                          |
|-----------------------------|--------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| <b>Communication style?</b> | (a) direct and concise ← RECOMMENDED<br>(b) direct, data-driven, comprehensive with tables<br>(c) conversational and collaborative<br>Press Enter to use recommended (a)                         |
| <b>Working pattern?</b>     | (a) iterative — start small, expand ambitiously ← RECOMMENDED<br>(b) plan-first — full specs before code<br>(c) prototype-fast — build quick, refine later<br>Press Enter to use recommended (a) |
| <b>AI delegation?</b>       | (a) AI does 70%, user guides 30% ← RECOMMENDED<br>(b) AI does 90%, user reviews 10%<br>(c) 50/50 collaboration<br>Press Enter to use recommended (a)                                             |

### Saved to:

```
~/portable-spec-kit/user-profile/user-profile-janesmith.md (global — asked once)
workspace/portable-spec-kit/user-profile/user-profile-janesmith.md (committed — persists)
```

**Personalized from first session.** Profile asked once, saved globally. On new projects: shown and confirmed — keep or customize per project. Customized profiles save to workspace. Press Enter through everything for recommended defaults.

## Your Projects

Project-specific details go in `agent/AGENT.md` (auto-created per project). This keeps the framework file portable while each project has its own stack, brand, and rules:

```
## Stack
| Layer | Technology |
|-----|-----|
| Frontend | Next.js + TypeScript |
| Backend | FastAPI |

## Brand
- Primary: #1B2A4A
- Accent: #4a6fa5
```

## Dev Stack Adaptation

The framework uses stack-agnostic language for testing, building, and linting. Each project's `agent/AGENT.md` defines the actual commands:

| Framework Says               | Project Defines                                                    |
|------------------------------|--------------------------------------------------------------------|
| "All tests passing"          | <code>npm test</code> / <code>pytest</code> / <code>go test</code> |
| "Type checking: zero errors" | <code>npx tsc --noEmit</code> / <code>mypy</code>                  |
| "Linting: zero errors"       | <code>eslint</code> / <code>ruff</code> / <code>golint</code>      |
| "Build succeeds"             | <code>npm run build</code> / <code>docker build</code>             |
| "PDF generation tool"        | WeasyPrint / Puppeteer / wkhtmltopdf                               |

**Stack-agnostic by design.** The framework never assumes a specific language, build tool, or test runner. It defines *what* to check — each project defines *how*.

# Sharing & Syncing the Framework

## Recommended: Dedicated GitHub Repo

### Single-Source-of-Truth Repo

One repo holds portable-spec-kit.md. All projects pull from it. Updates pushed back from any project.

```
# Setup (once)
mkdir portable-spec-kit && cd portable-spec-kit
git init && cp /path/to/portable-spec-kit.md .
git add . && git commit -m "Framework v1.0"
gh repo create portable-spec-kit --private --push

# Pull into any project
curl -s0 https://raw.githubusercontent.com/you/portable-spec-kit/main/portable-spec-kit.md

# Push updates back from any project
cp portable-spec-kit.md ~/portable-spec-kit/
cd ~/portable-spec-kit && git add . && git commit -m "Updated from ProjectX" && git push
```

✓ Version controlled ✓ Pull anywhere ✓ Change history ✓ Cross-machine sync

## Other Options

| Method               | How                                                             | Best For                        |
|----------------------|-----------------------------------------------------------------|---------------------------------|
| <b>Git Submodule</b> | <code>git submodule add repo .framework</code>                  | Auto-sync across repos          |
| <b>GitHub Gist</b>   | <code>gh gist create portable-spec-kit.md</code>                | Quick sharing, browser editable |
| <b>Shell Alias</b>   | <code>alias init-dev='curl -s0 URL/portable-spec-kit.md'</code> | One-command setup               |



# Update Workflow Across Projects

---



## Portable vs Project-Specific

---

| In portable-spec-kit.md (Portable Framework) | In agent/AGENT.md (Per Project) |
|----------------------------------------------|---------------------------------|
| Git/push/commit rules                        | Tech stack choices              |
| Security practices                           | Brand colors & fonts            |
| Testing methodology                          | AI model/provider config        |
| Naming conventions                           | Dev server port                 |
| Project templates & auto-scan                | Deployment specifics            |
| Code review standards                        | Project-specific rules          |
| Spec-driven development flow                 | API endpoints & data model      |
| Version numbering scheme                     | Database schema                 |

# Examples Included

The repo includes two example projects showing the framework in action:

| Example               | What It Shows                                                                               | Best For                                        |
|-----------------------|---------------------------------------------------------------------------------------------|-------------------------------------------------|
| examples/<br>starter/ | Fresh project right after setup. Self-documenting README explains every file and directory. | Understanding the framework — <b>start here</b> |
| examples/my-app/      | Mid-development Next.js + Supabase project. 11/16 tasks done, 24 tests at 92% coverage.     | Seeing a realistic project in progress          |

## Starter Project Structure

```
examples/starter/
├─ portable-spec-kit.md      ← Framework file (source)
├─ .portable-spec-kit/      ← Kit config (user profiles)
│   └─ user-profile/
├─ CLAUDE.md                ← Symlink (Claude Code)
├─ .cursorrules              ← Symlink (Cursor)
├─ WORKSPACE_CONTEXT.md     ← Auto-created workspace state
├─ README.md                 ← Self-documenting (explains everything)
├─ agent/
│   └─ AGENT.md              ← Stack: TBD (waiting for your specs)
│   └─ AGENT_CONTEXT.md     ← Status: "Setup – waiting for specs"
│   └─ SPECS.md              ← Empty – ready for requirements
│   └─ PLANS.md              ← Empty – ready for architecture & methodology
│   └─ TASKS.md              ← 1/5 tasks (project initialized)
│   └─ RELEASES.md          ← v0.1 placeholder
```

## Mid-Development Project

```
examples/my-app/
├─ agent/
│   └─ AGENT.md              ← Next.js + Supabase configured
│   └─ AGENT_CONTEXT.md     ← v0.1: 11/16 tasks, 24 tests
│   └─ SPECS.md              ← 8 features, priorities, criteria
│   └─ PLANS.md              ← Data model, APIs, 3 phases
│   └─ TASKS.md              ← 5 modules with progress
│   └─ RELEASES.md          ← Categorized changelog
```

**One file. Any project. Your standards.**

The framework travels with you — consistent engineering practices across every project, every team, every AI session.

Portable Spec Kit

# Thank You

**Author:** Dr. Aqib Mumtaz, Ph.D.

**Specialization:** Ph.D. Computer Science — Artificial  
Intelligence

**Research:** Multimodal AI, Healthcare AI, Autonomous  
Surveillance

**Email:** [aqib.mumtaz@gmail.com](mailto:aqib.mumtaz@gmail.com)

**LinkedIn:** [linkedin.com/in/aqibmumtaz](https://linkedin.com/in/aqibmumtaz)

**GitHub:** [github.com/aqibmumtaz](https://github.com/aqibmumtaz)

v0.6.98 • June 2026

© 2026 Dr. Aqib Mumtaz. All rights reserved.